

This is a story about
our participation in
one of our customer's
project

Scala SWAT

TACKLING A 1 BILLION MEMBER SOCIAL NETWORK



Artur Bańkowski
artur@evojam.com
[@abankowski](https://twitter.com/abankowski)



Previous Experience

Datasets I have previously worked on were much smaller!

~10 millions records

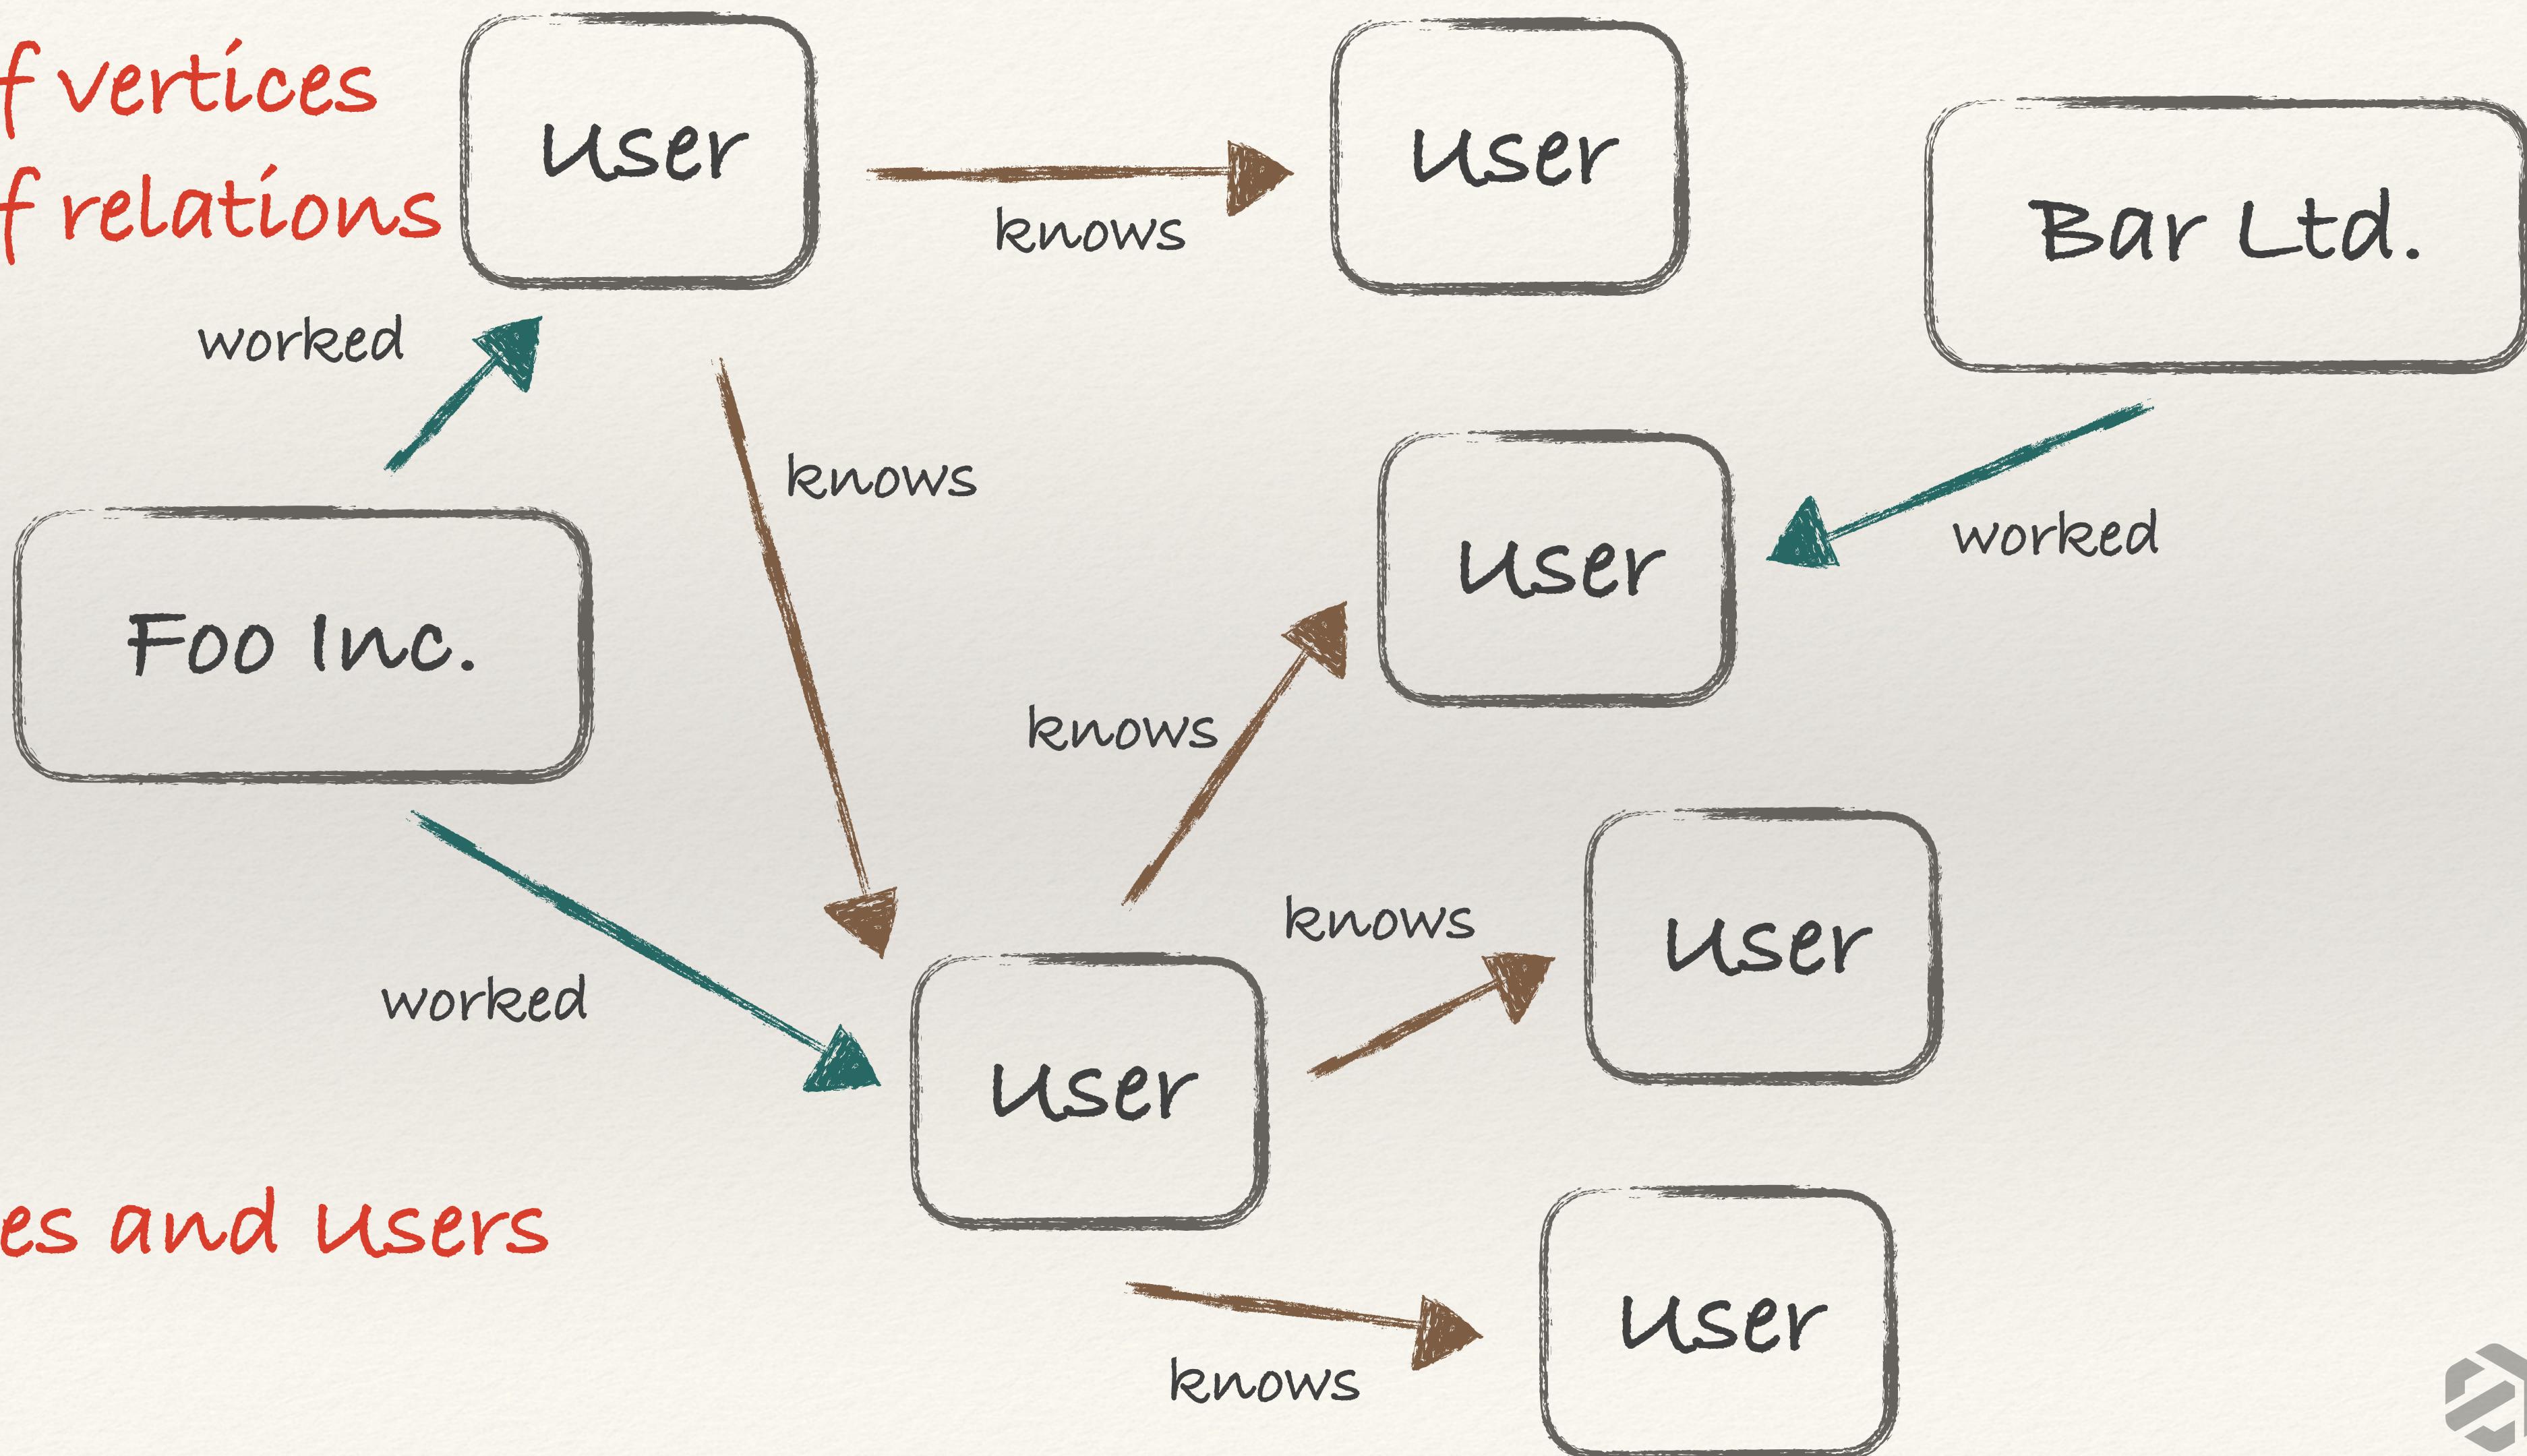
~60 millions documents

~60 million vertices graph (non-commercial)

This time we were about to deal with a much bigger social network, which can easily be represented as a graph. We were going to join the team...

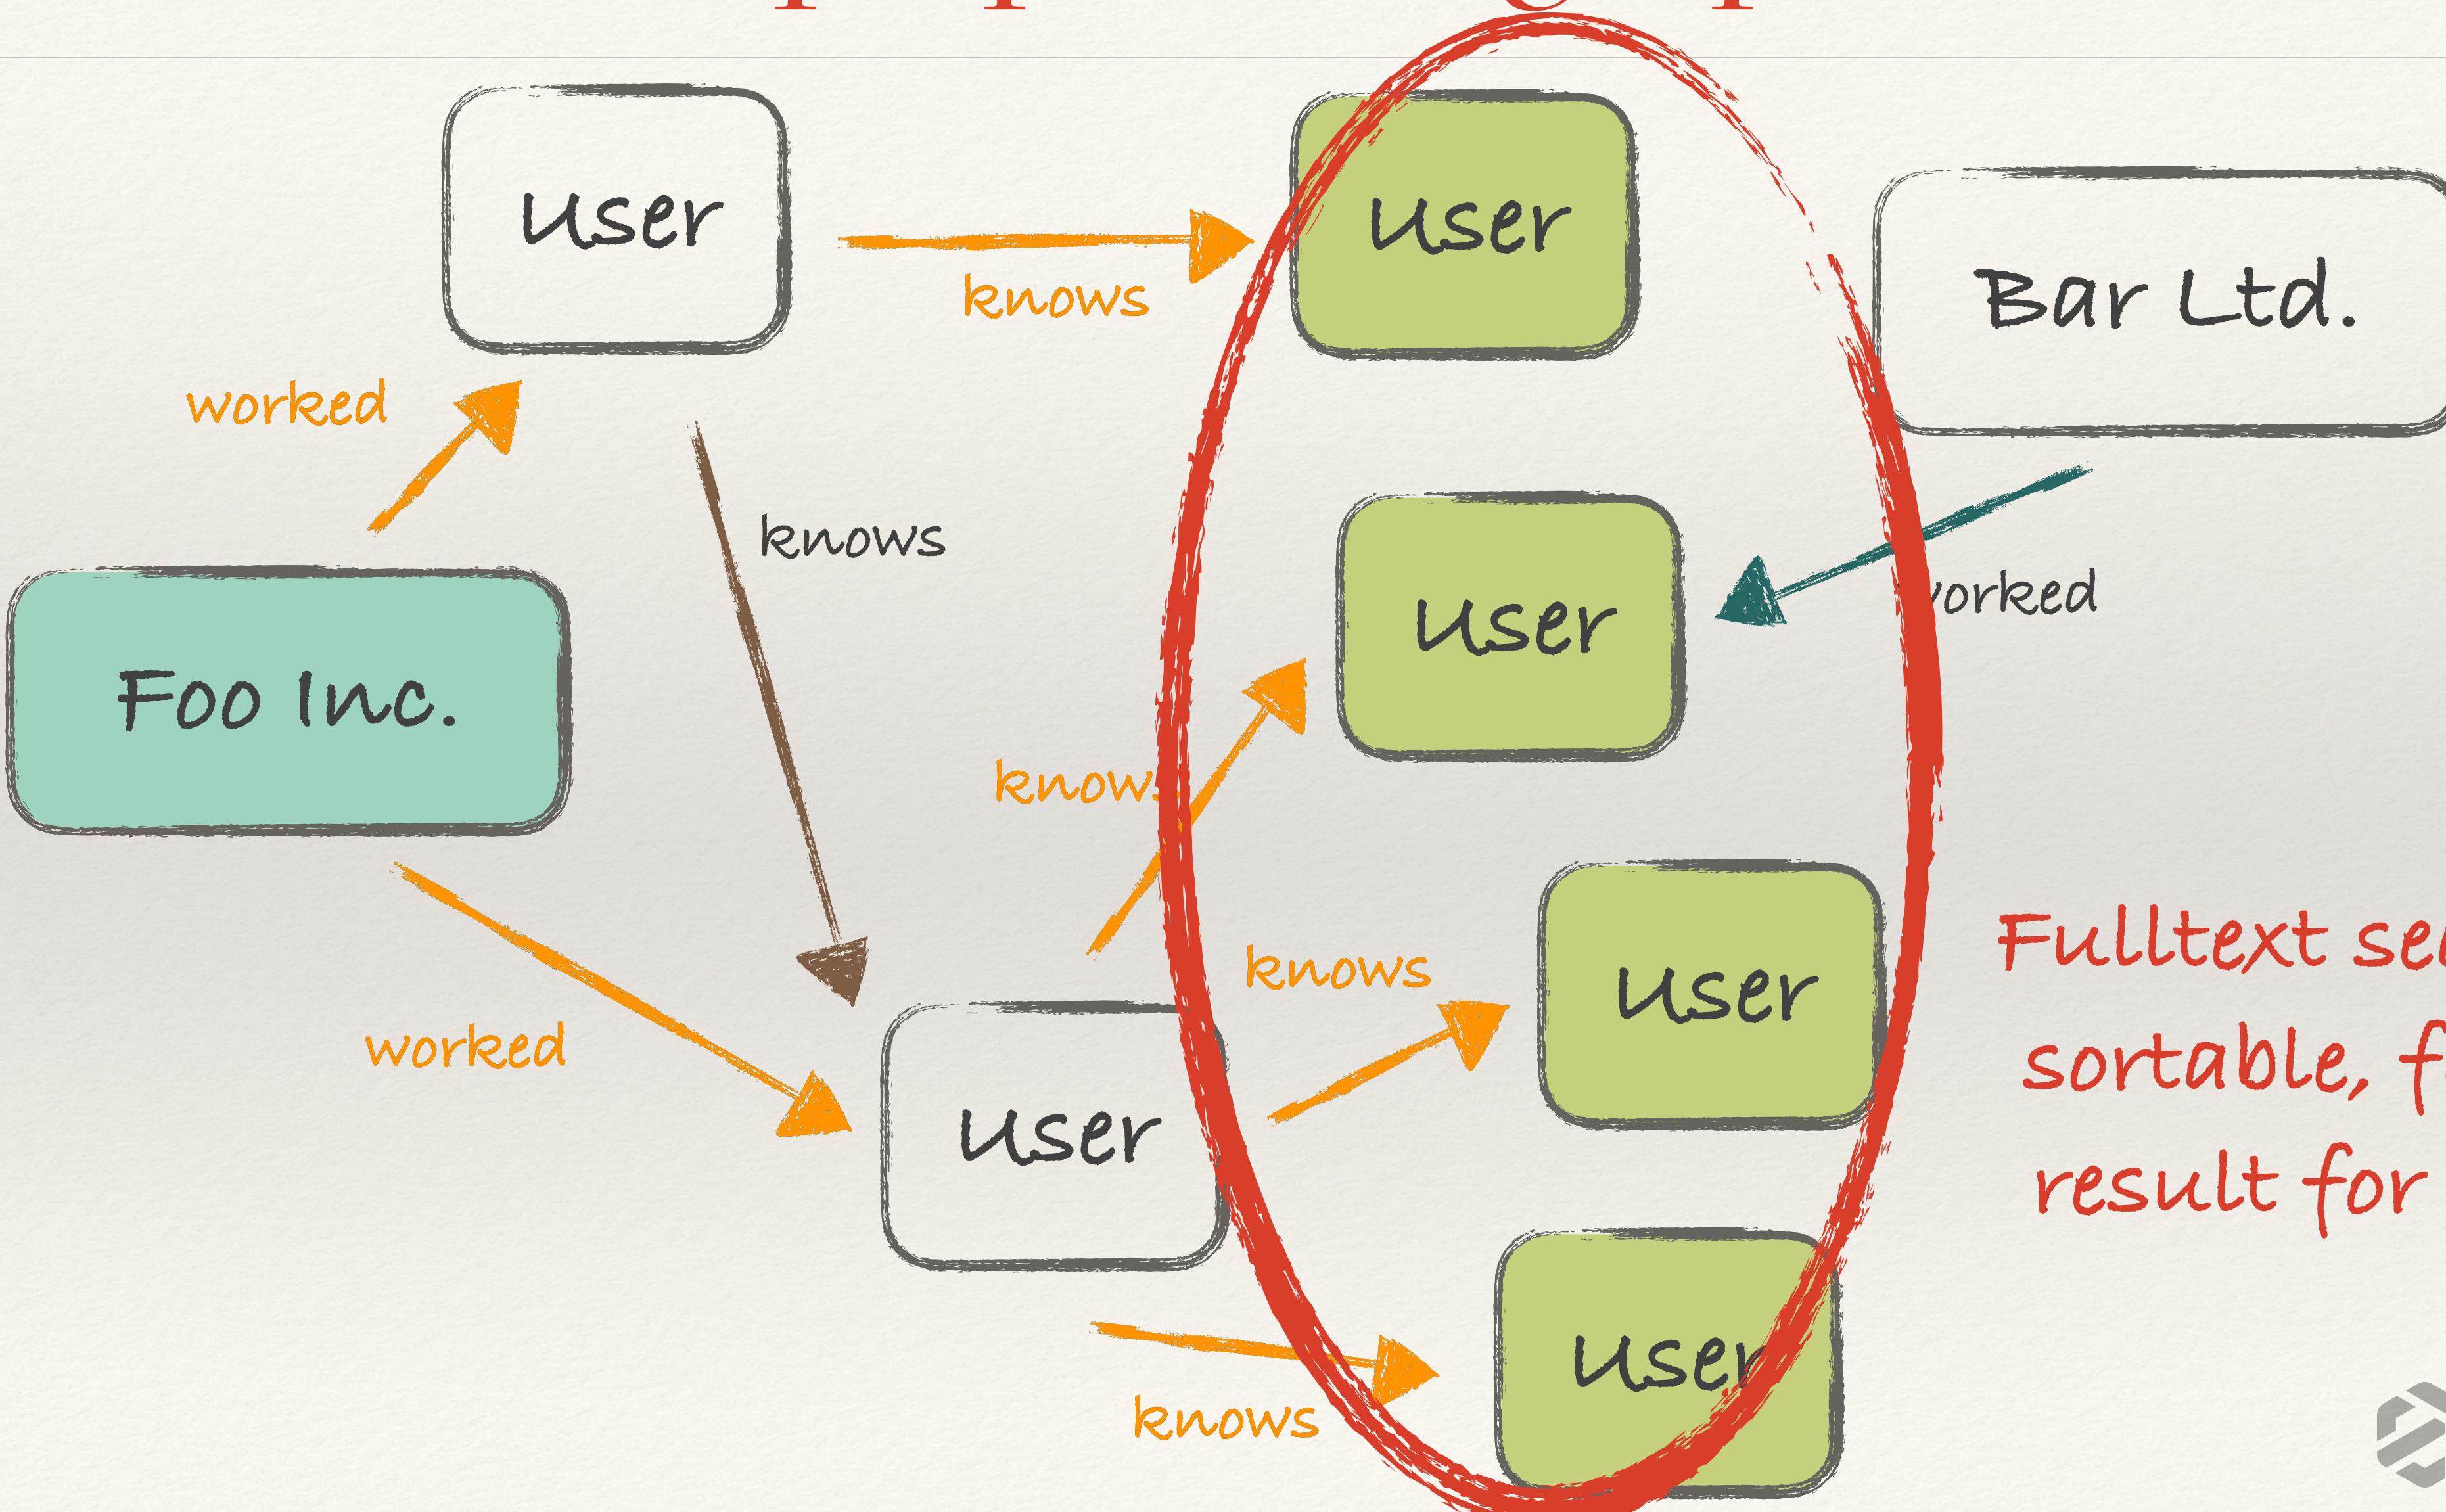
Graph structure

2 types of vertices
2 types of relations



Companies and users

The concept: provide graph subset



835,272,759

~~1 billion challenge~~

When we finally put
our hands on the data,
we have realized that
datasize is smaller
than expected

835,272,759 vertices

751,857,081 users

83,415,678 companies

6,956,990,209 relations

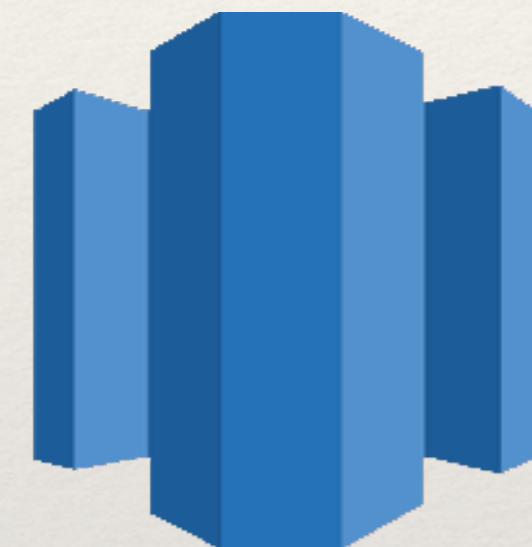
What more have we
found?

Subsets from thousands to millions users

Existing app workflow

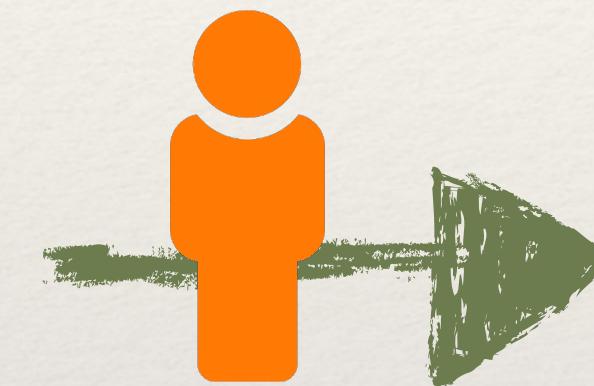
This was already used by final customers

extract subset
eg: 3m profiles



Amazon **Redshift**

750m profiles



JSON files

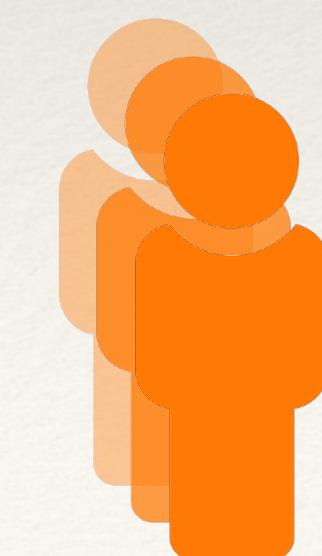
data duplication



elastic

only manually selected
60m incl. duplicates

*Manual process,
takes few days from
user perspective*



One Engineer-Evening
Multiple custom scripts

PoC - The Goal

Handle *1 billion* profiles automatically

Our primary goal

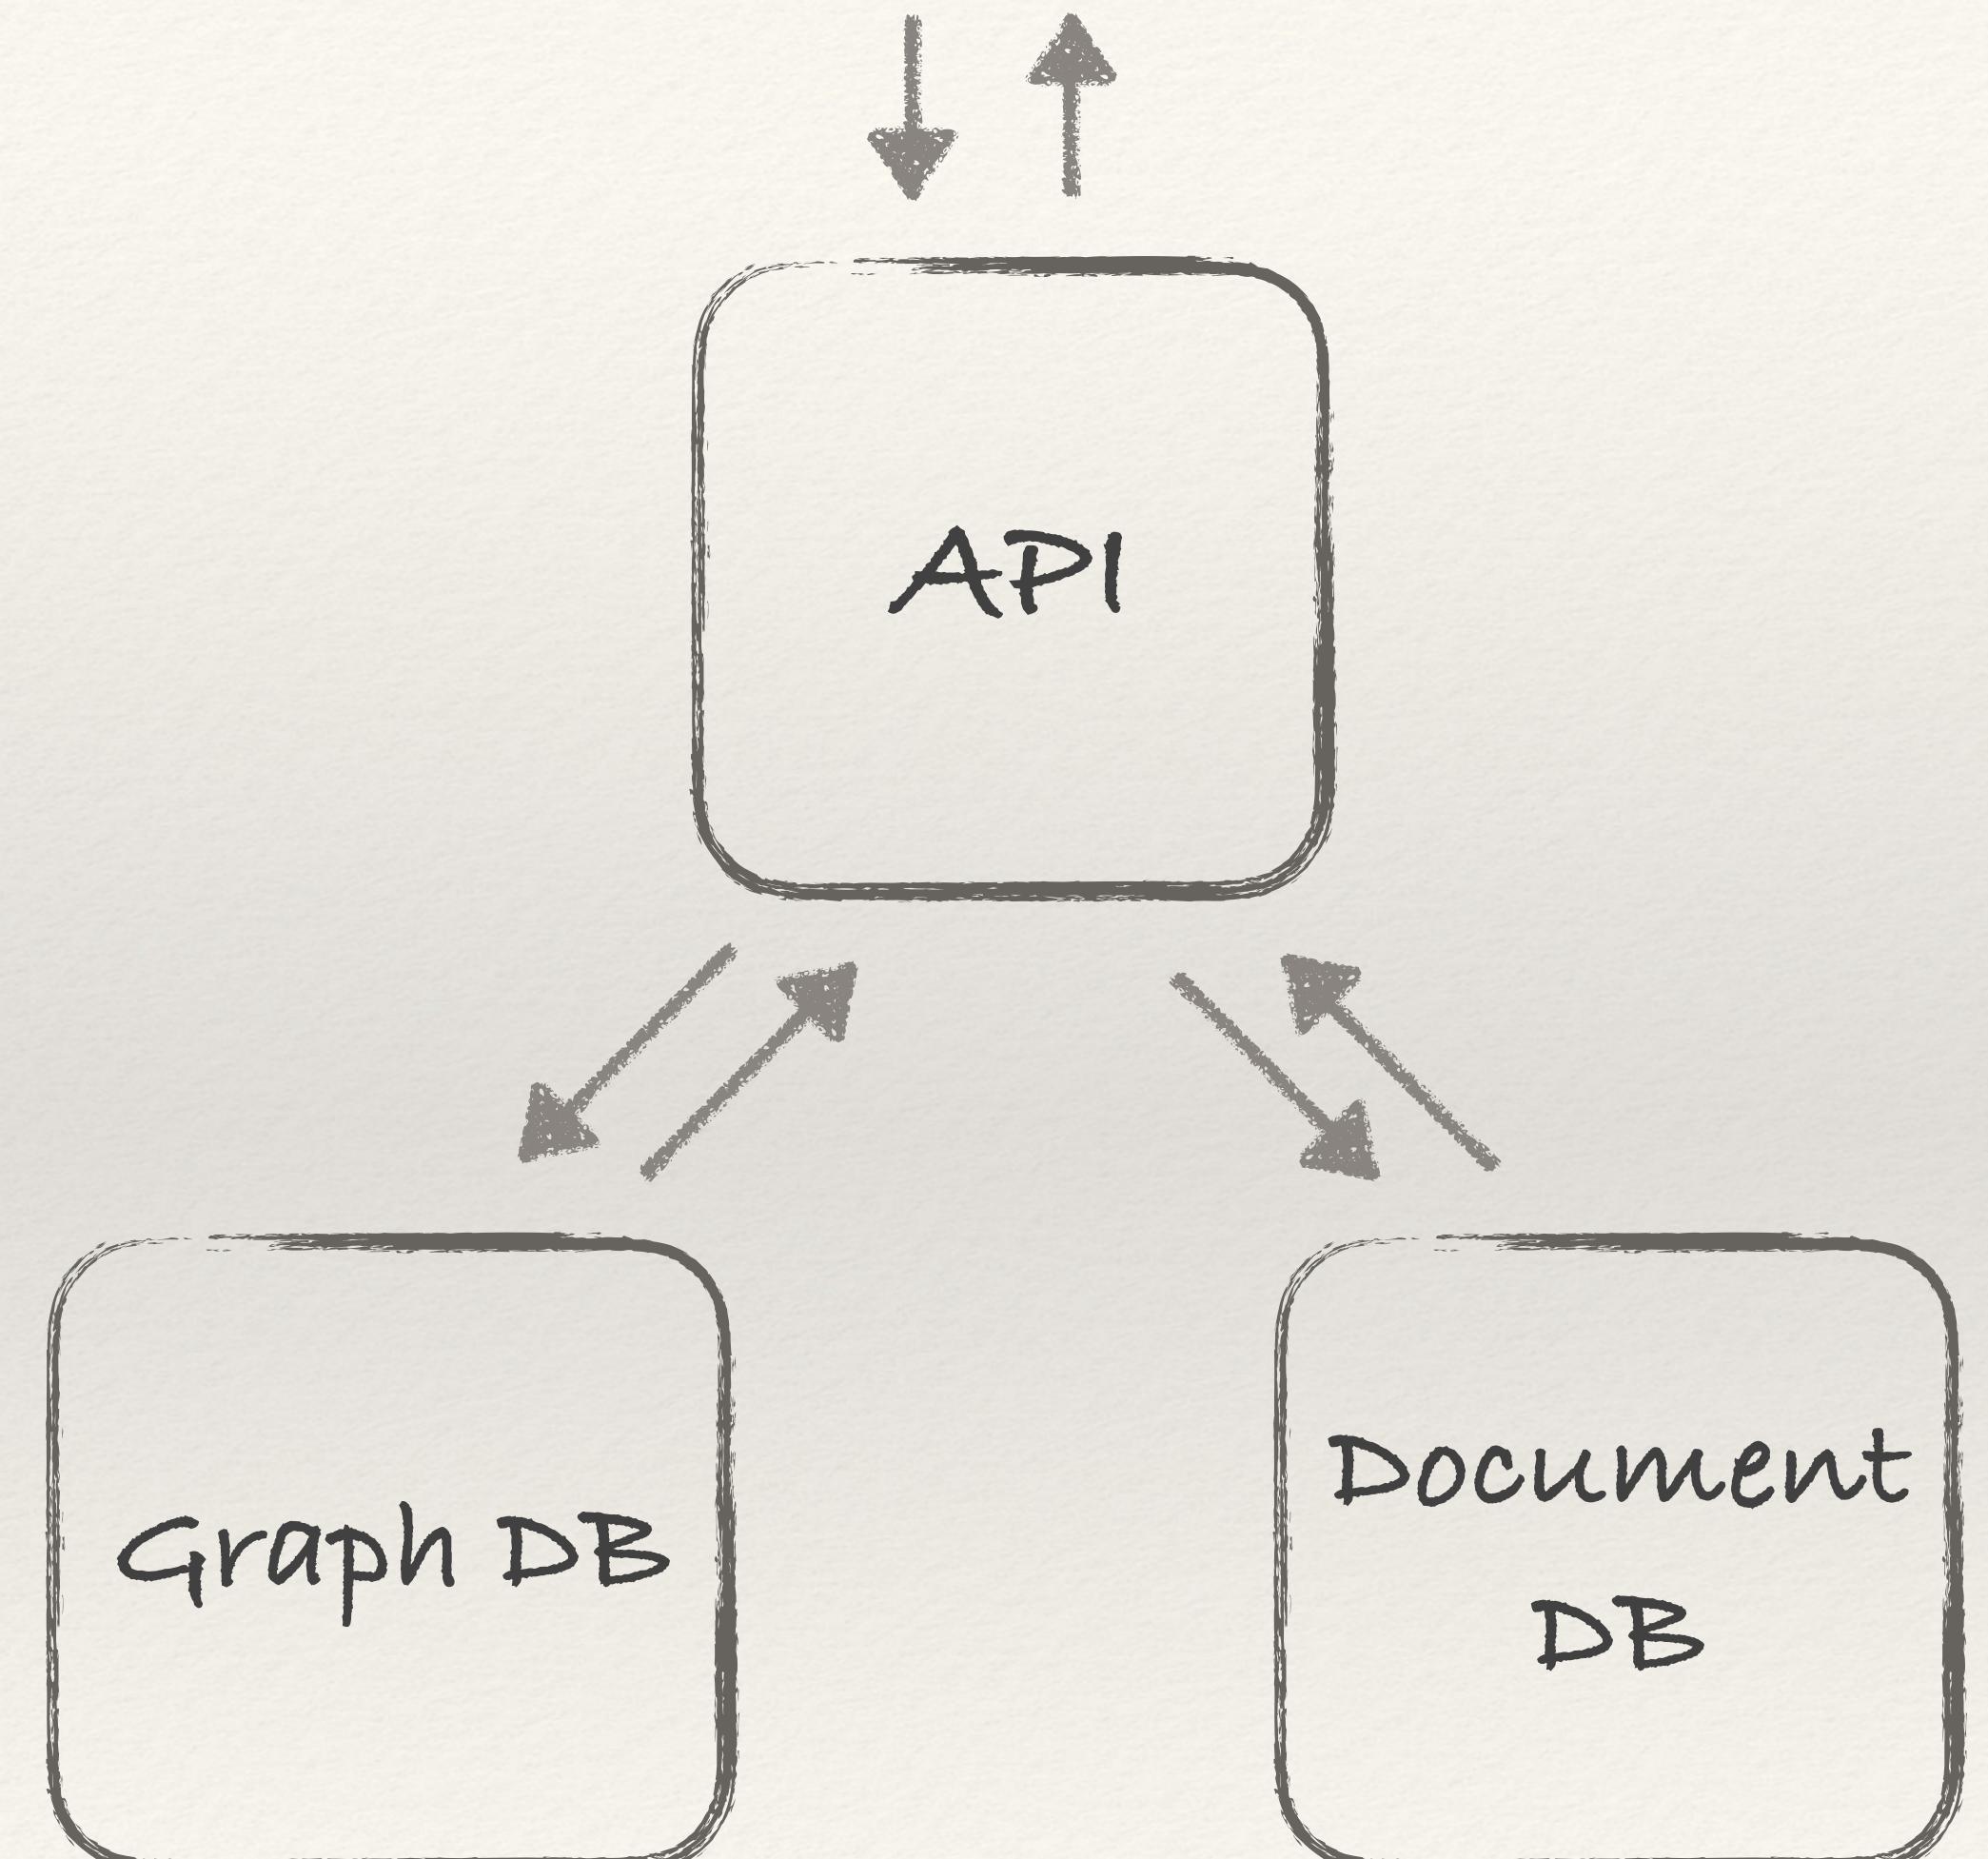
PoC - Definition of done

REST API with search

- Graph traversable on demand
- First results available under 1 minute
- Entire subset ready in few minutes

PoC - Concept

Whole dataset
stored in two
engines



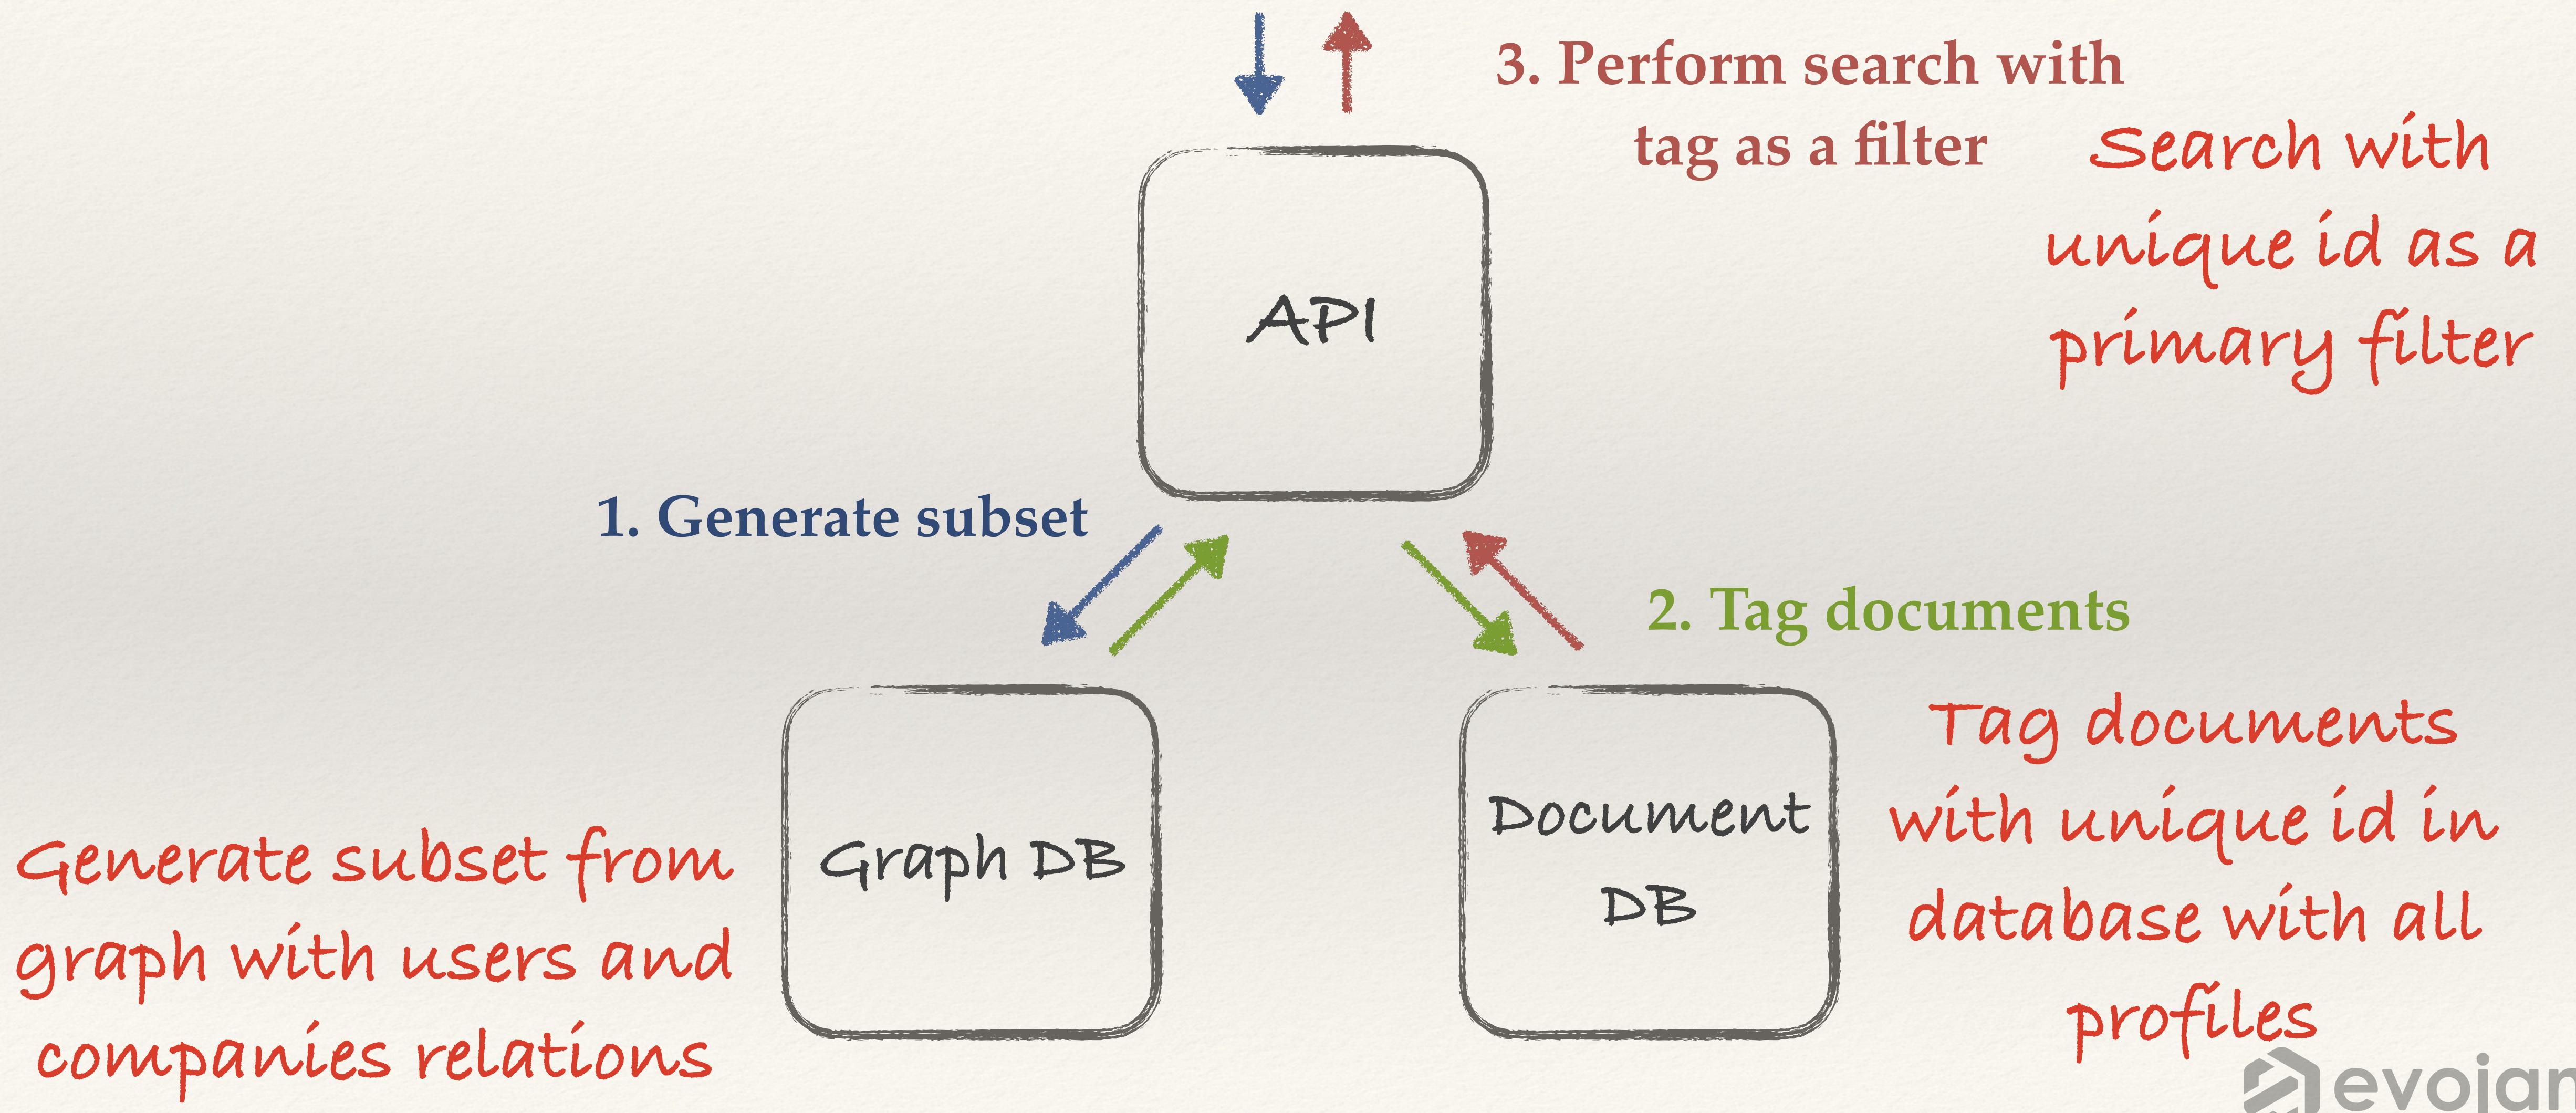
All relations
6,956,990,209 relations

All profiles
751,857,081 profiles

Why graph DB?

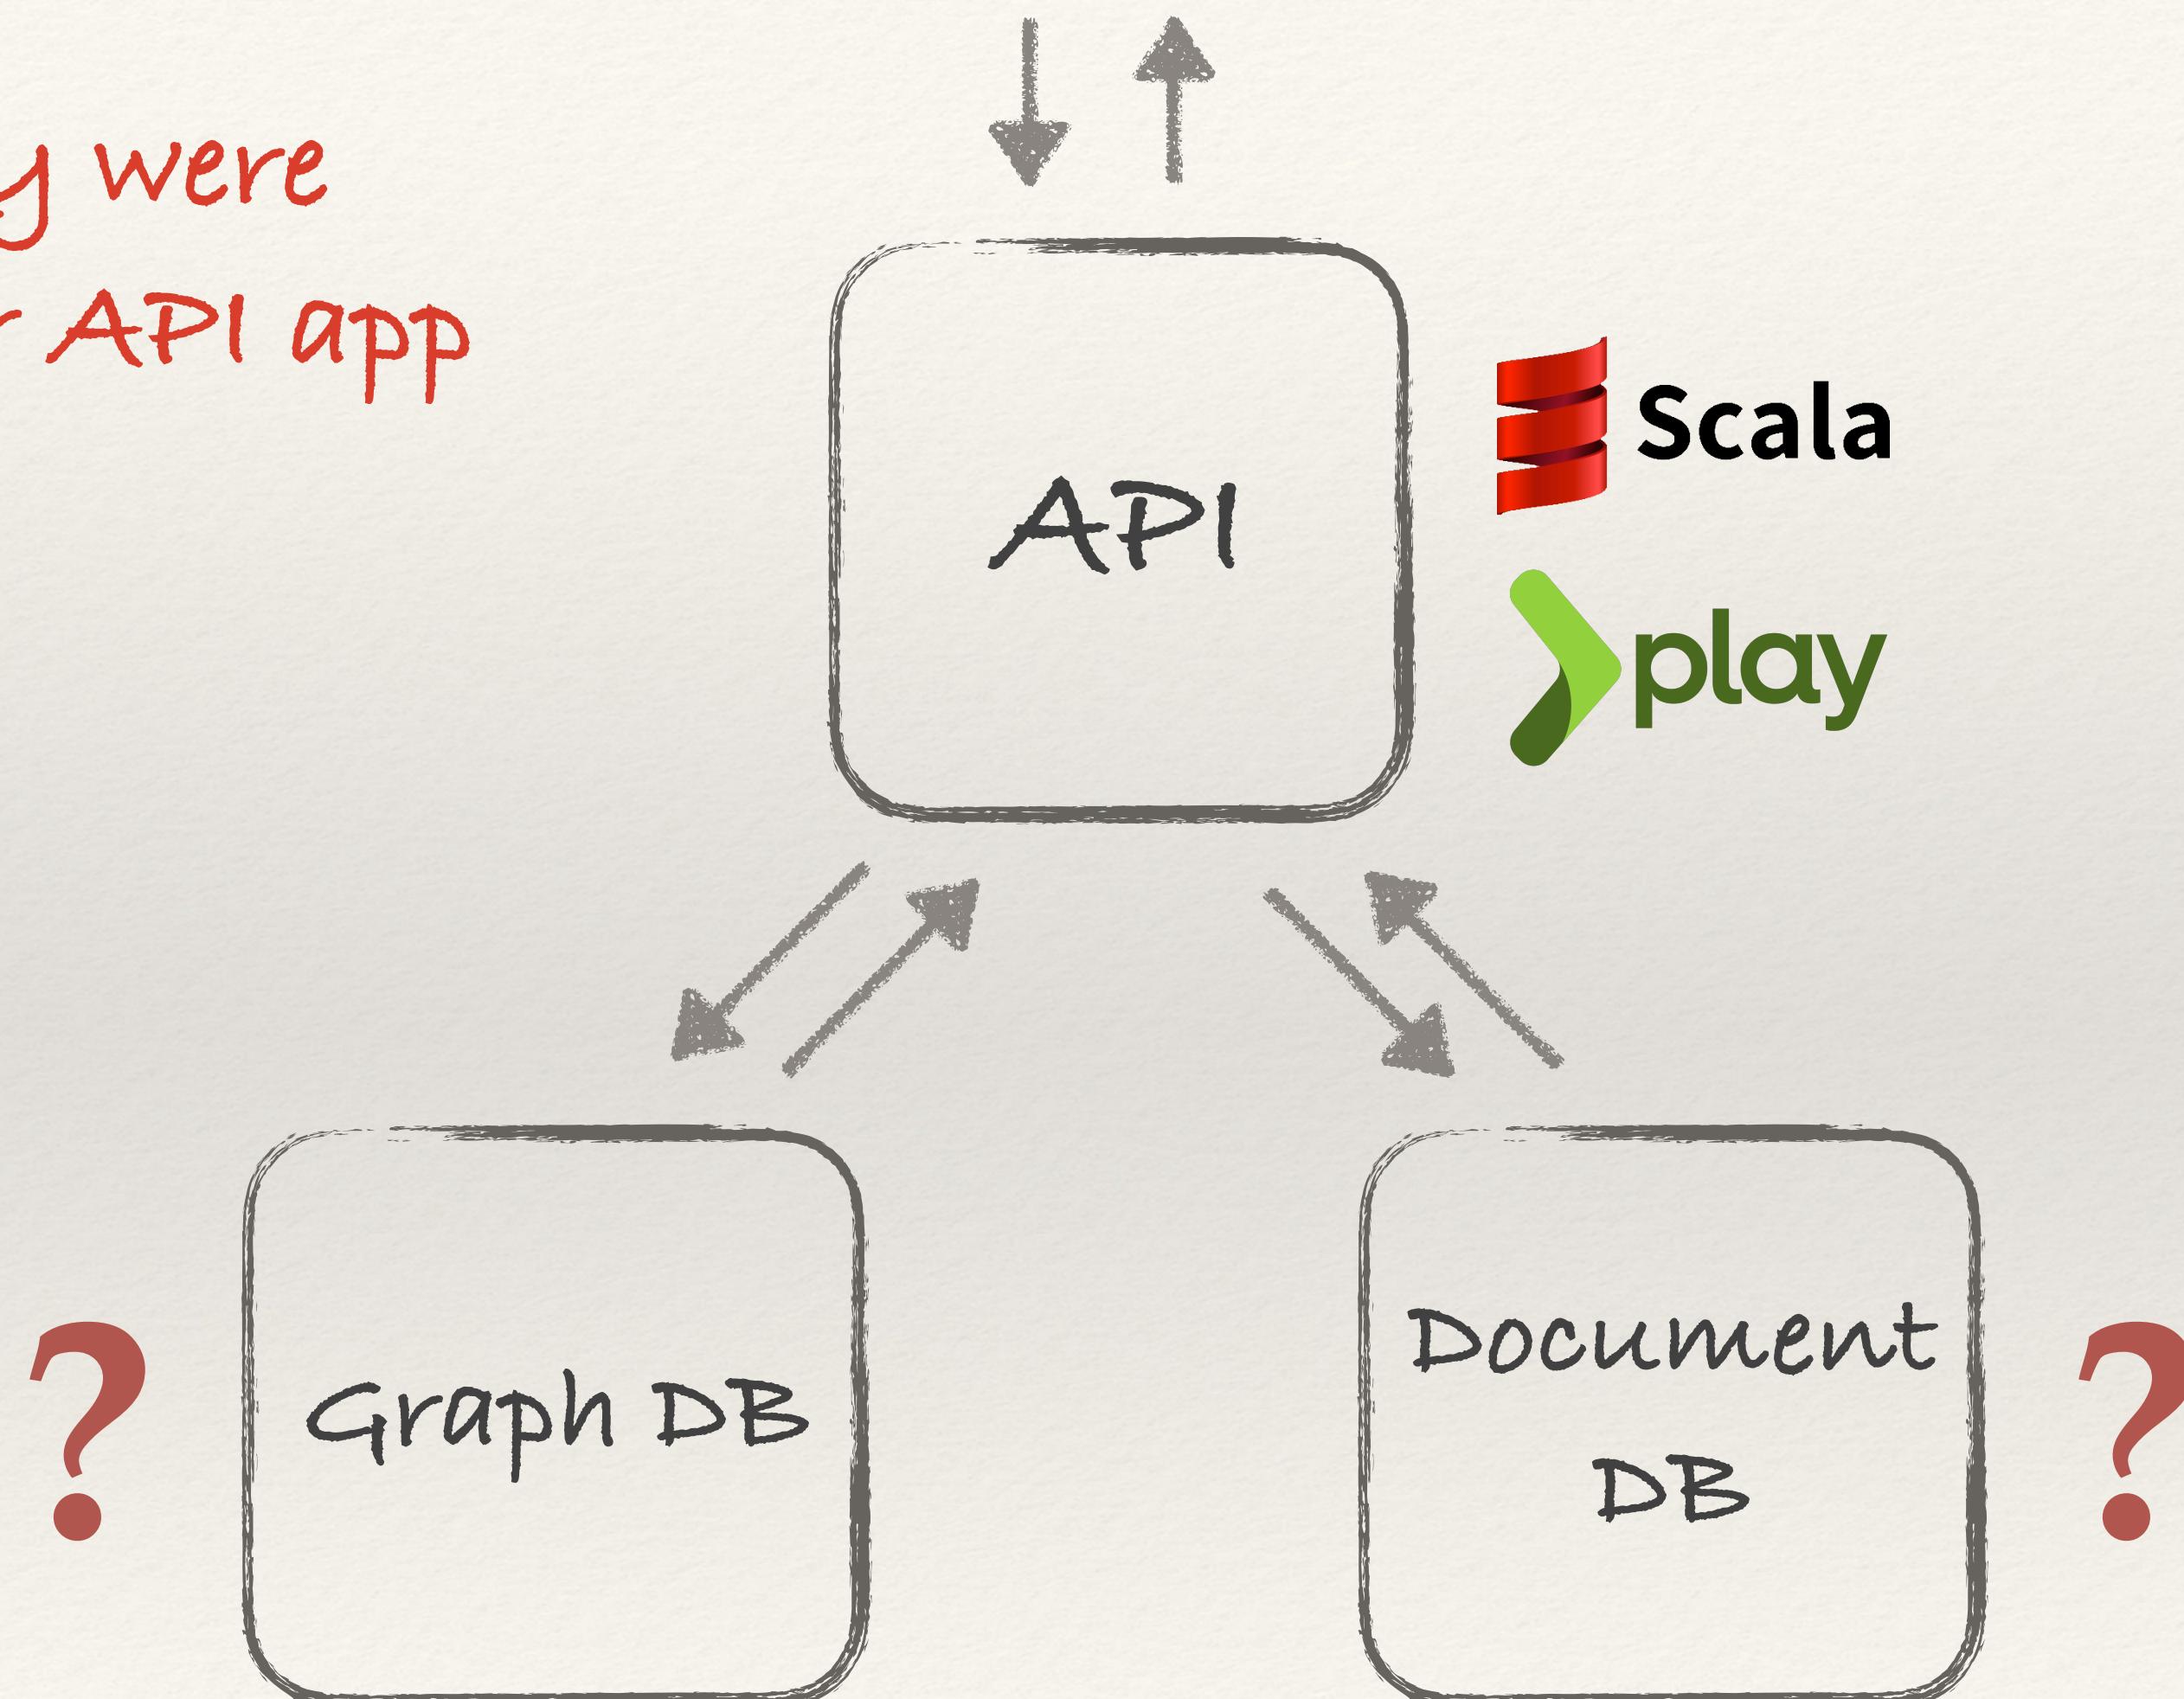
- Fast graph traversal
- Easily extendable with new relations and vertices
- Convenient algorithm description

PoC - Flow



Weapon of choice

Scala and Play were natural choice for API app



Some research required for databases

First Steps: Extraction

To make a research we had to put our hands on the real data



Amazon **Redshift**

Extract anonymized
profiles, companies
and relations

Cleanup data, sort
and generate input
files

*few days to pull, streaming
with Akka and Slick*

First Steps: Pushing forward



Push profiles for
searching purposes

two tools,
highly dependent on db engines



Push vertices and
relations for traversal

Fulltext Searchable Document DB



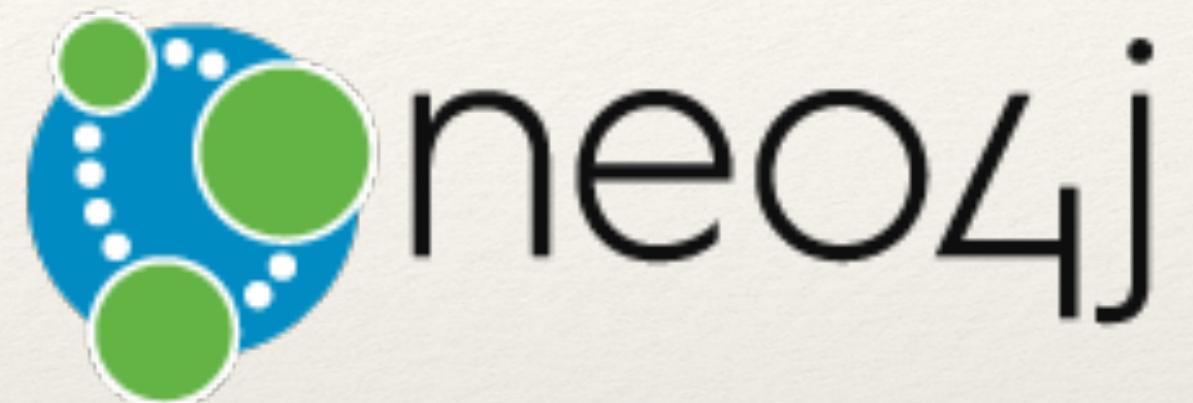
- Mature
- Horizontally scalable
- Fast indexing (~3k documents per second on the single node)
- Well documented
- With scala libraries:
 - <https://github.com/sksamuel/elastic4s>
 - <https://github.com/evojam/play-elastic4s>

we already had significant experience with scaling Elastic Search
for 80 millions

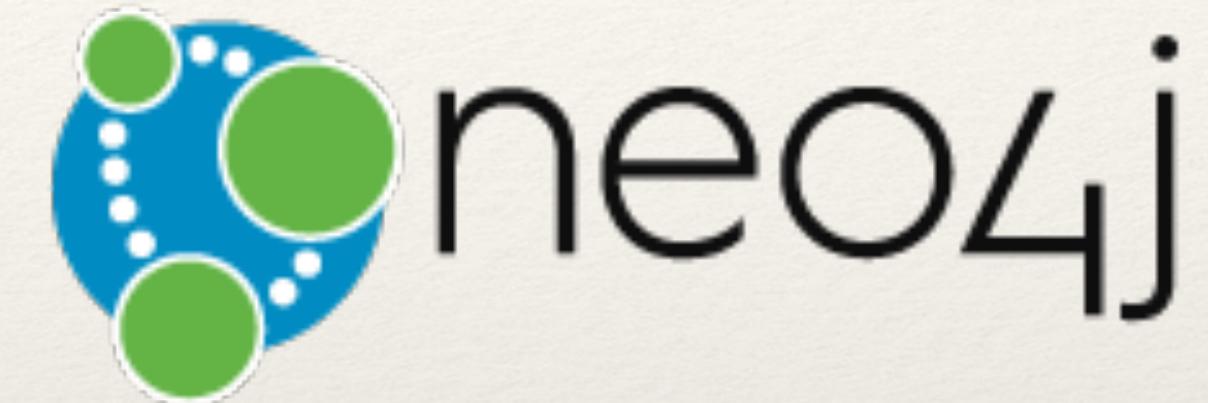
Which graph DB?



Considered three engines,
OrientDB and Neo4j in
community editions



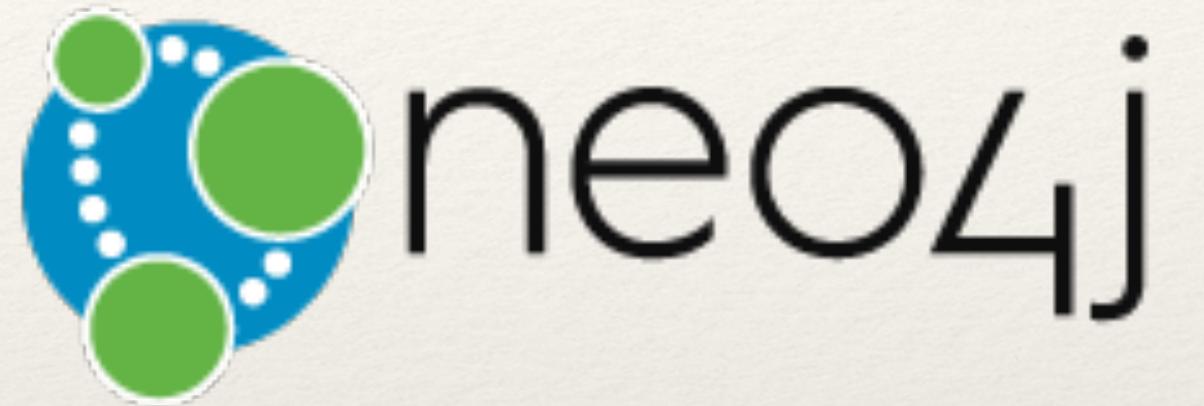
Goodbye Titan



- Performed well in ~60M tests
- fast traversing
- Development has been put on hold?



- No batch insertion, slow initial load
- Stalled writes after 200 millions of vertices with relations
- Horror stories in the internet



Neo4j FTW

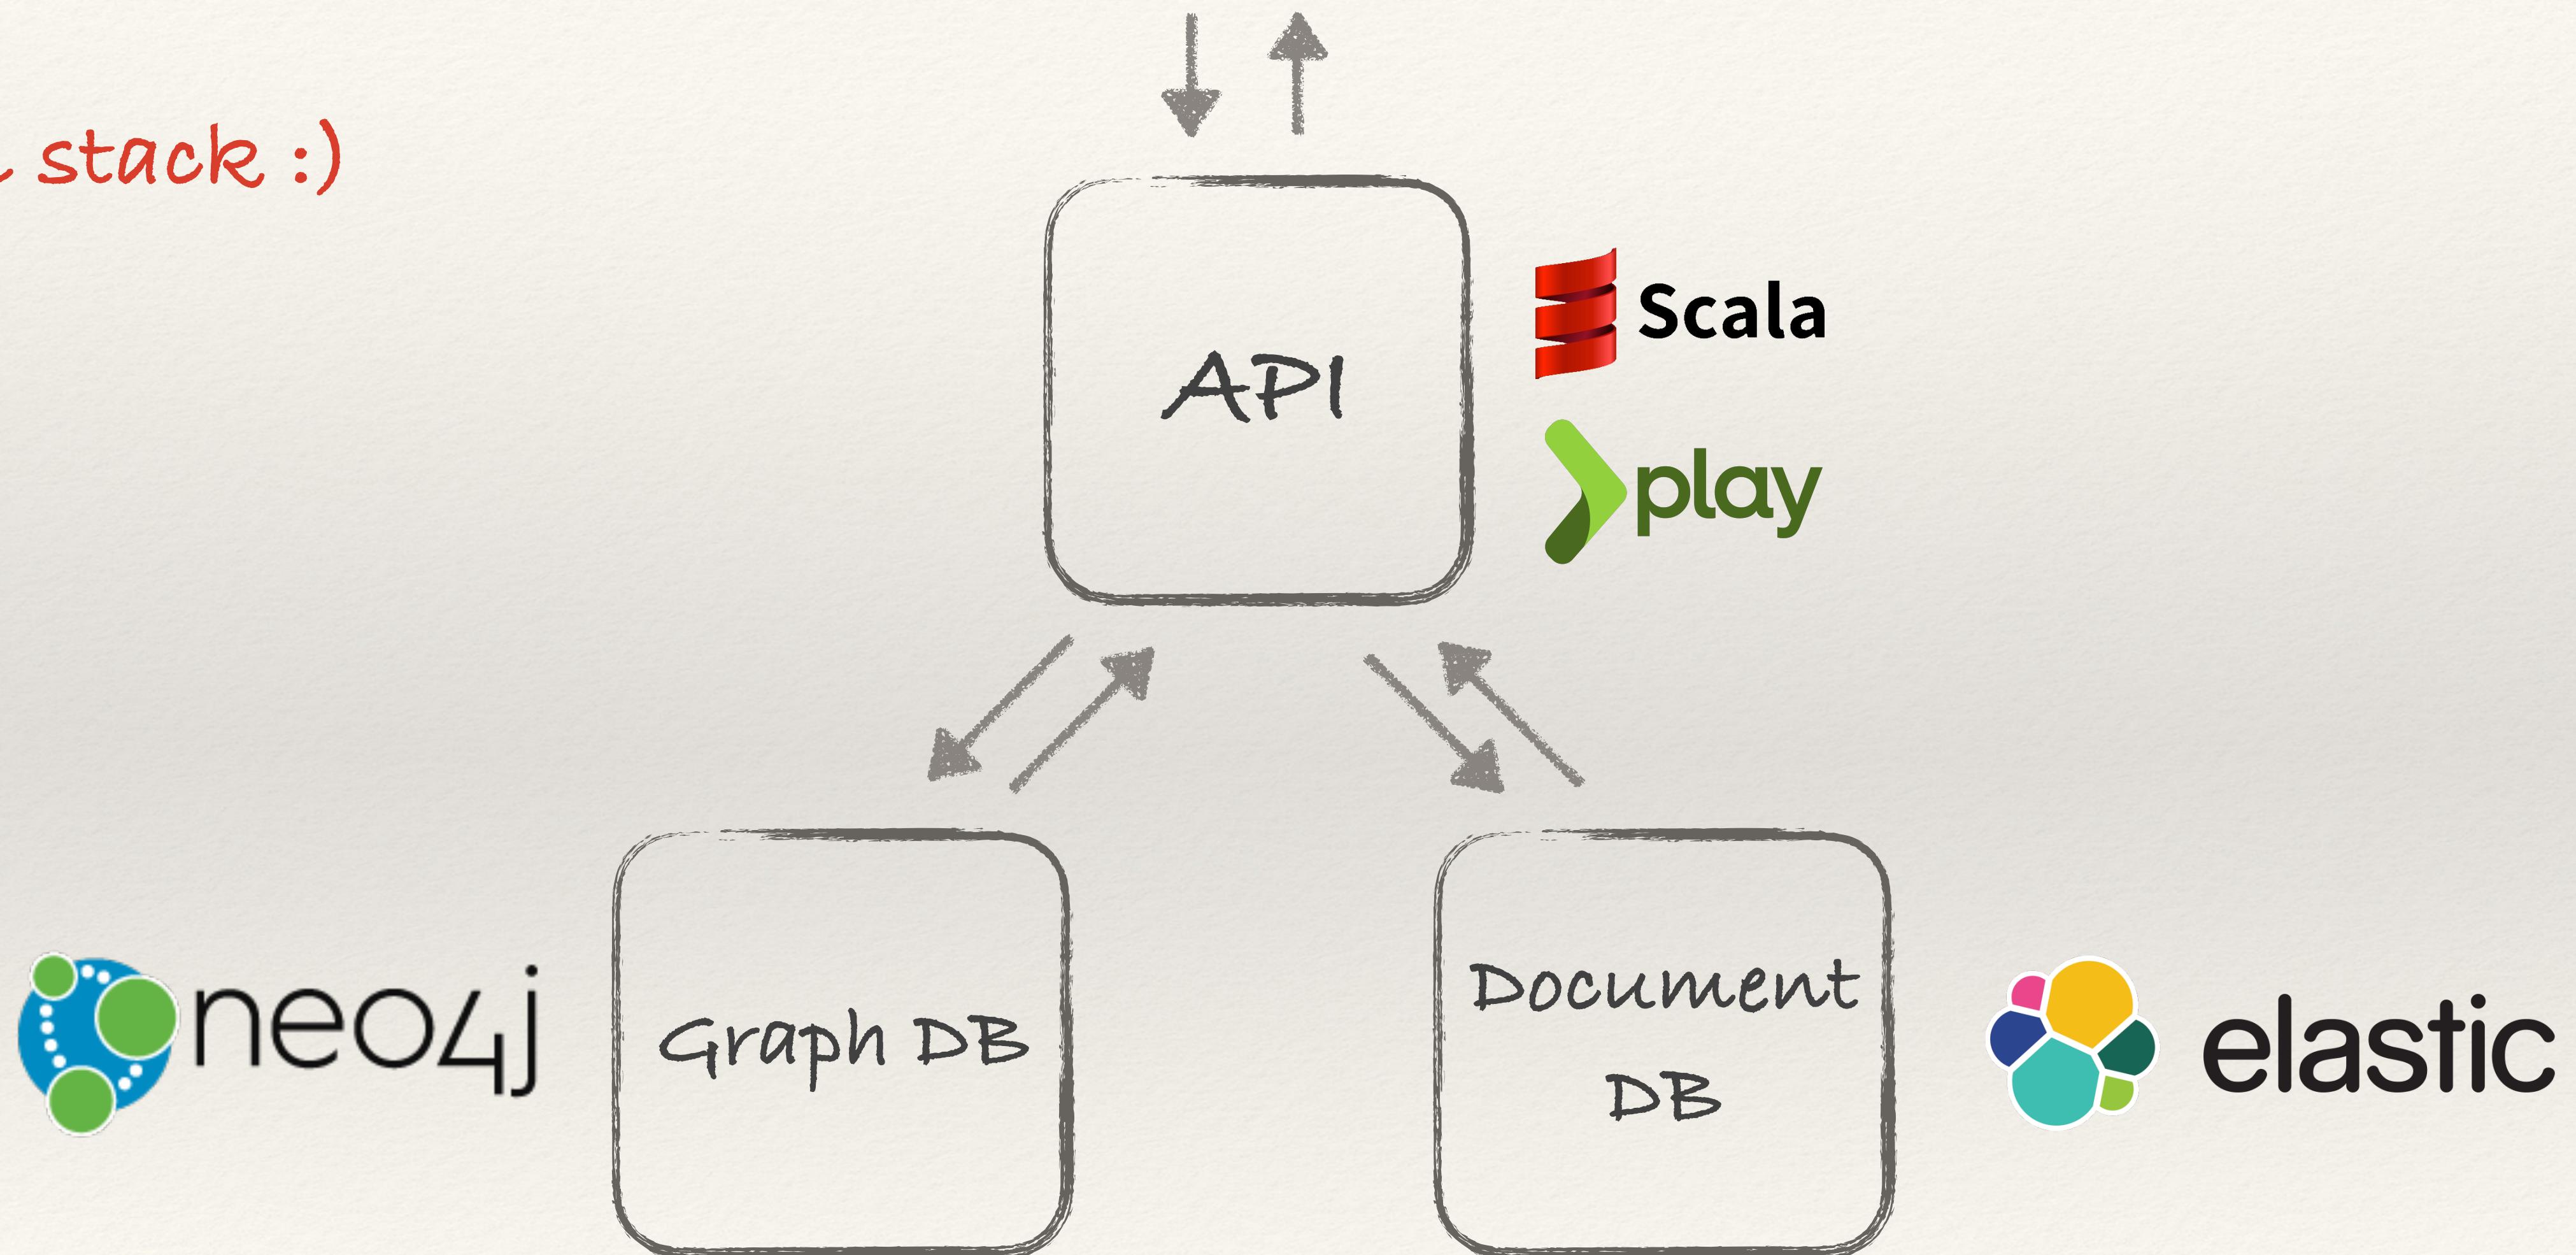
<https://github.com/AnormCypher/AnormCypher>



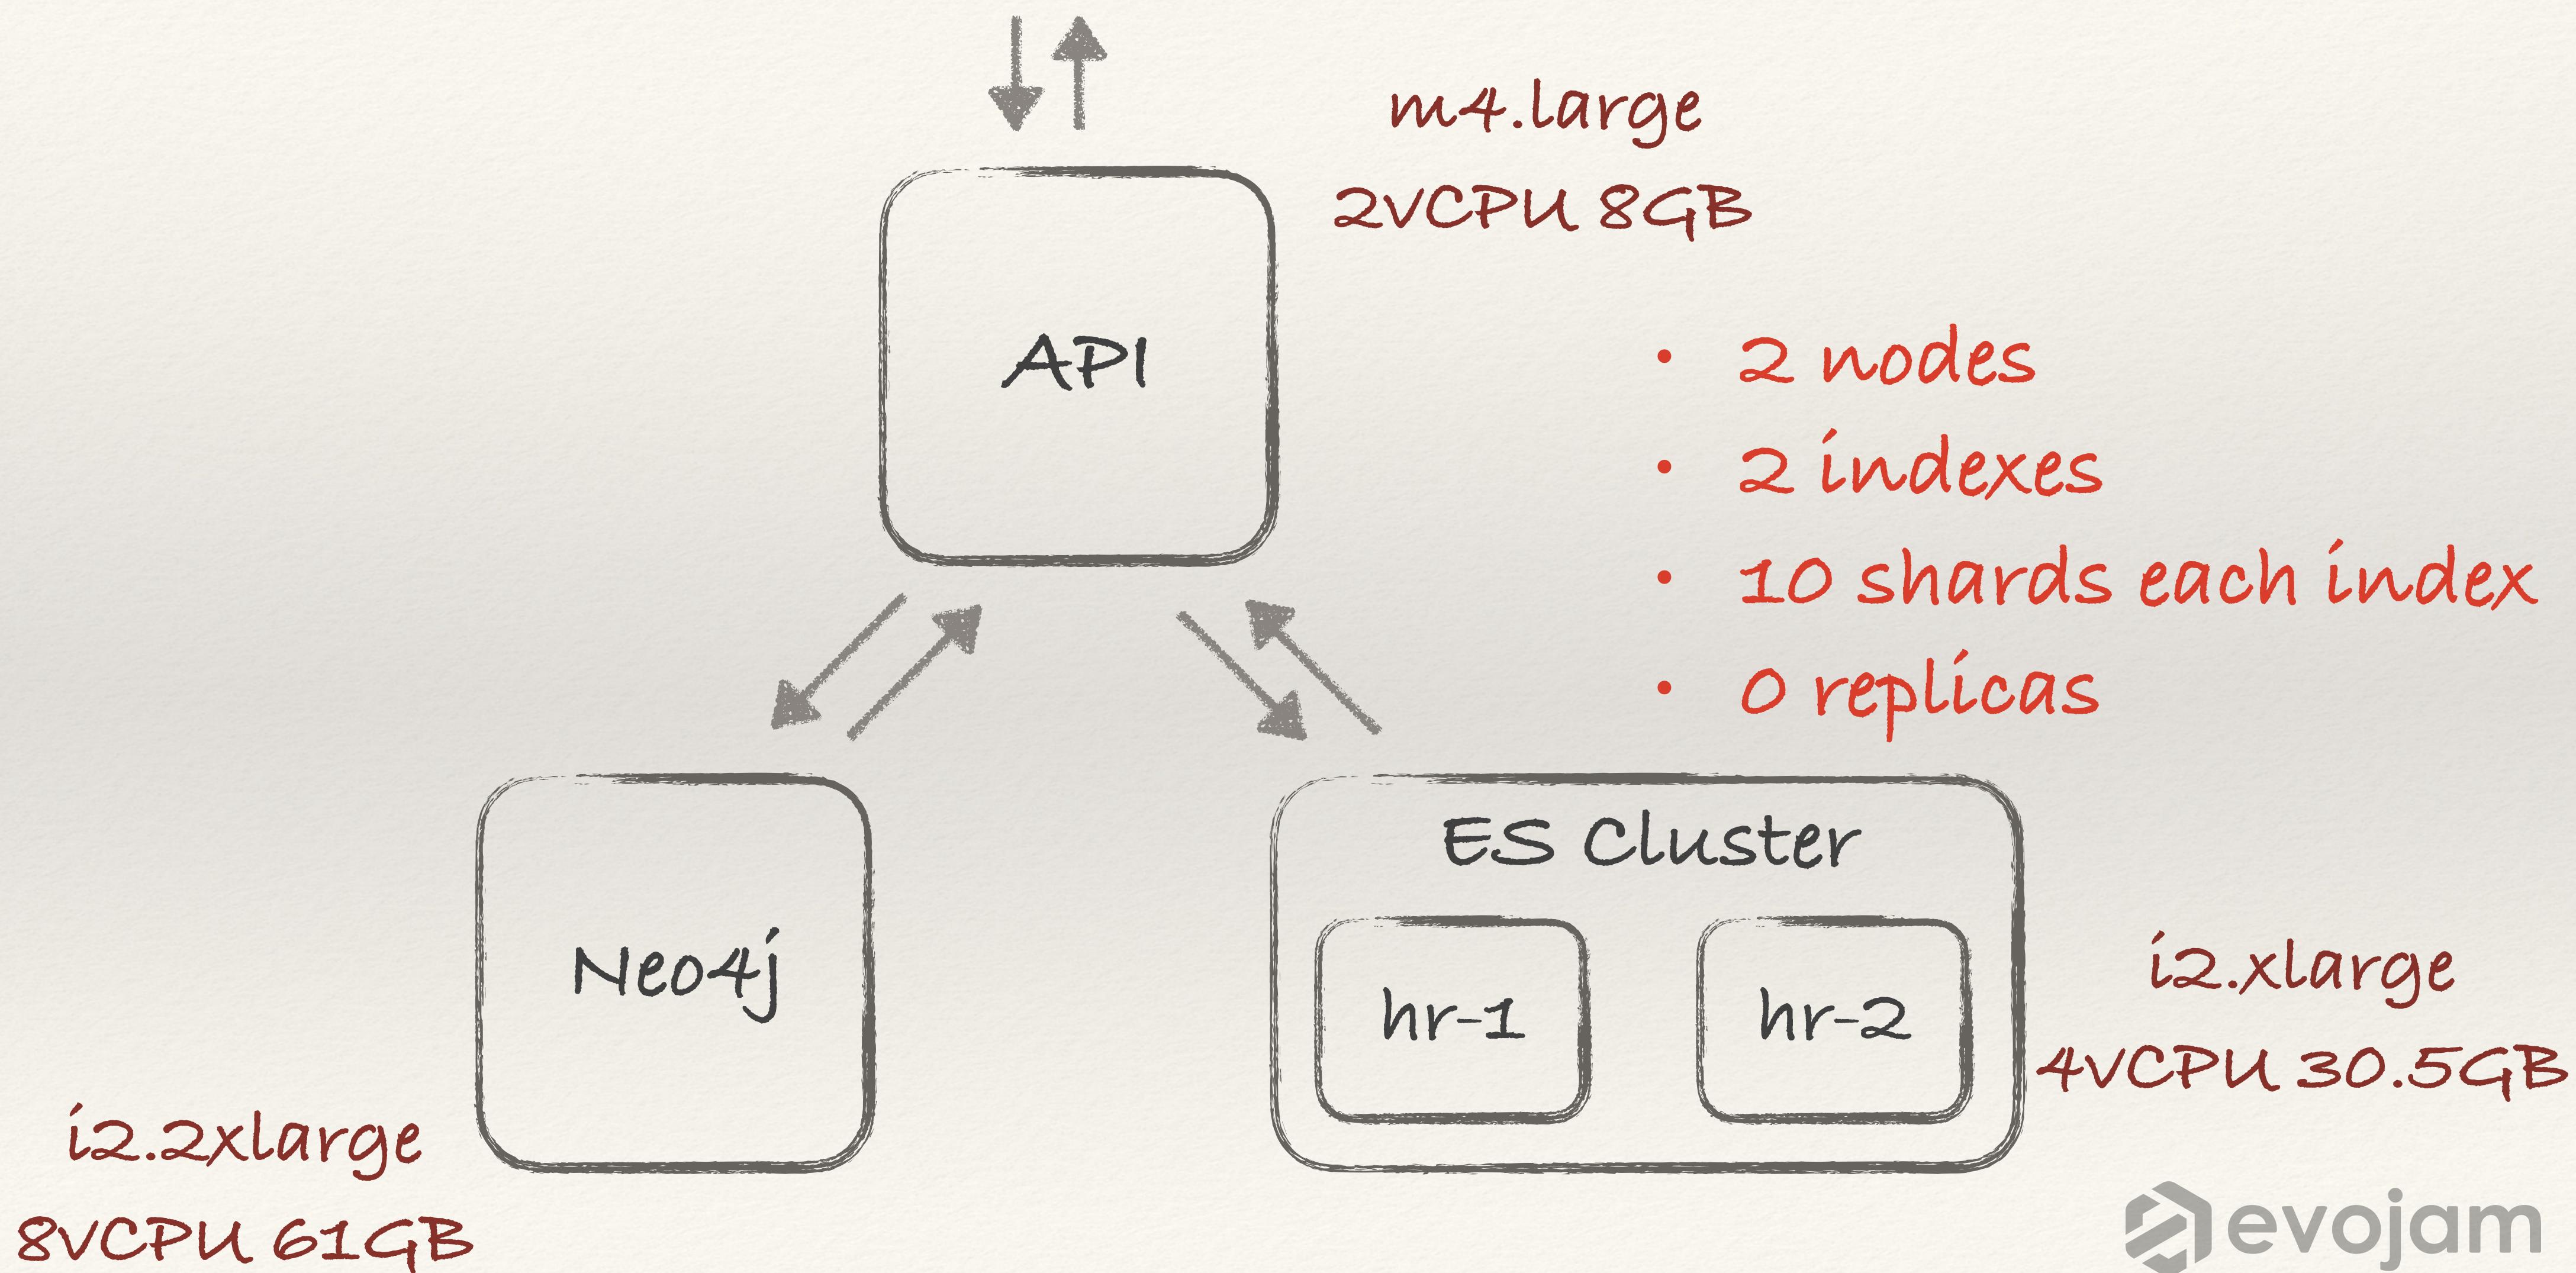
- Fast
- Already known
- Convenient in Scala with AnormCypher
- Offline bulk loading
- Result streaming

Weapon of choice

Final stack :)



Final Setup on AWS



Step #1 - Bulk loading into Neo4j

```
Importing the contents of these files into data/graph.db.
```

```
[...]
```

```
IMPORT DONE in 3h 24m 58s 140ms. Imported:  
835273352 nodes  
6956990209 relationships  
0 properties
```

It took 12 hours on 2 times smaller Amazon instance

Step #2 - Bulk loading into ES



Akka advantage:
CPU utilization,
parallel data
enrichment,
human readable

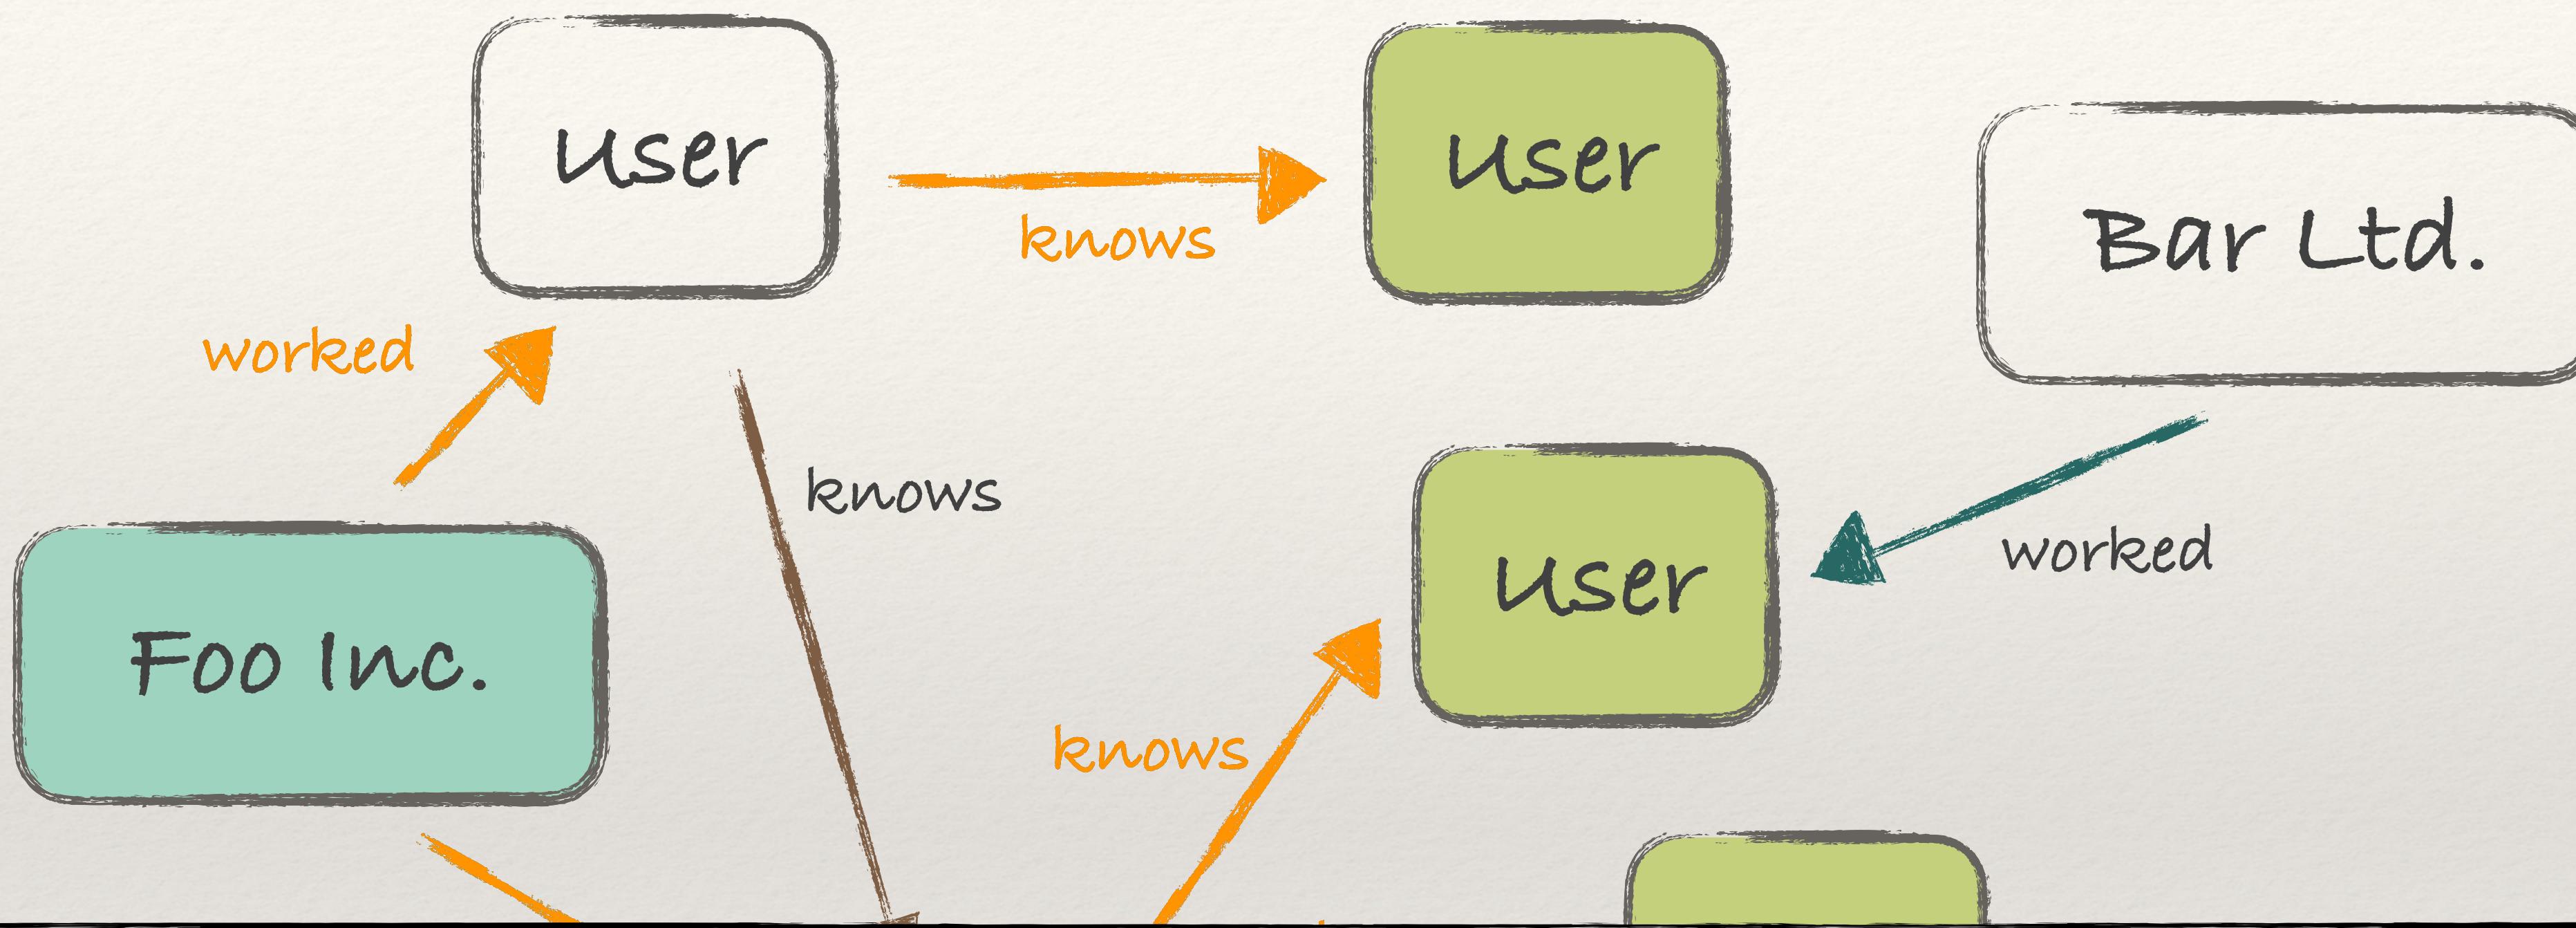
1. Create Source from CSV file, frame by \n
2. Decode id from ByteString, generate user json
3. Group by the bulk size (eg.: 4500)
4. Throttle
5. Execute bulk insert into the ElasticSearch

Step #2 - Bulk loading into ES

Flow description
with Akka
Streams

```
FileIO.fromFile(new File(sourceOfUserIds))
    .via(
        Framing.delimiter(ByteString('\n'),
            maximumFrameLength = 1024,
            allowTruncation = false))
    .mapAsyncUnordered(16)(prepareUserJson)
    .grouped(4500)
    .throttle(
        elements = 2,
        per = 2 seconds,
        maximumBurst = 2,
        mode = ThrottleMode.Shaping)
    .mapAsyncUnordered(2)(executeBulkInsert)
    .runWith(Sink.ignore)
```

Step #3 - Tagging



```
MATCH (c:Company)-[:worked]-(employees:User)-[:knows]-(acquaintance:User)  
WHERE ID(c)={foo-inc-id} AND NOT (acquaintance)-[:worked]-(c)  
RETURN DISTINCT ID(acquaintance)
```

Neo4j traversal query in CYpher

Akka Streams to the rescue

```
val idEnum : Enumeratee[CypherRow] = _  
  
val src =  
  Source.fromPublisher(Streams.enumeratorToPublisher(idEnum))  
    .map(_.data.head.asInstanceOf[BigDecimal].toInt)  
    .via(new TimeoutOnEmptyBuffer())  
    .map(UserId(_))  
    .mapAsyncUnordered(parallelism = 1)(id =>  
      dao.tagProfiles(id, companyId))
```

Readable flow

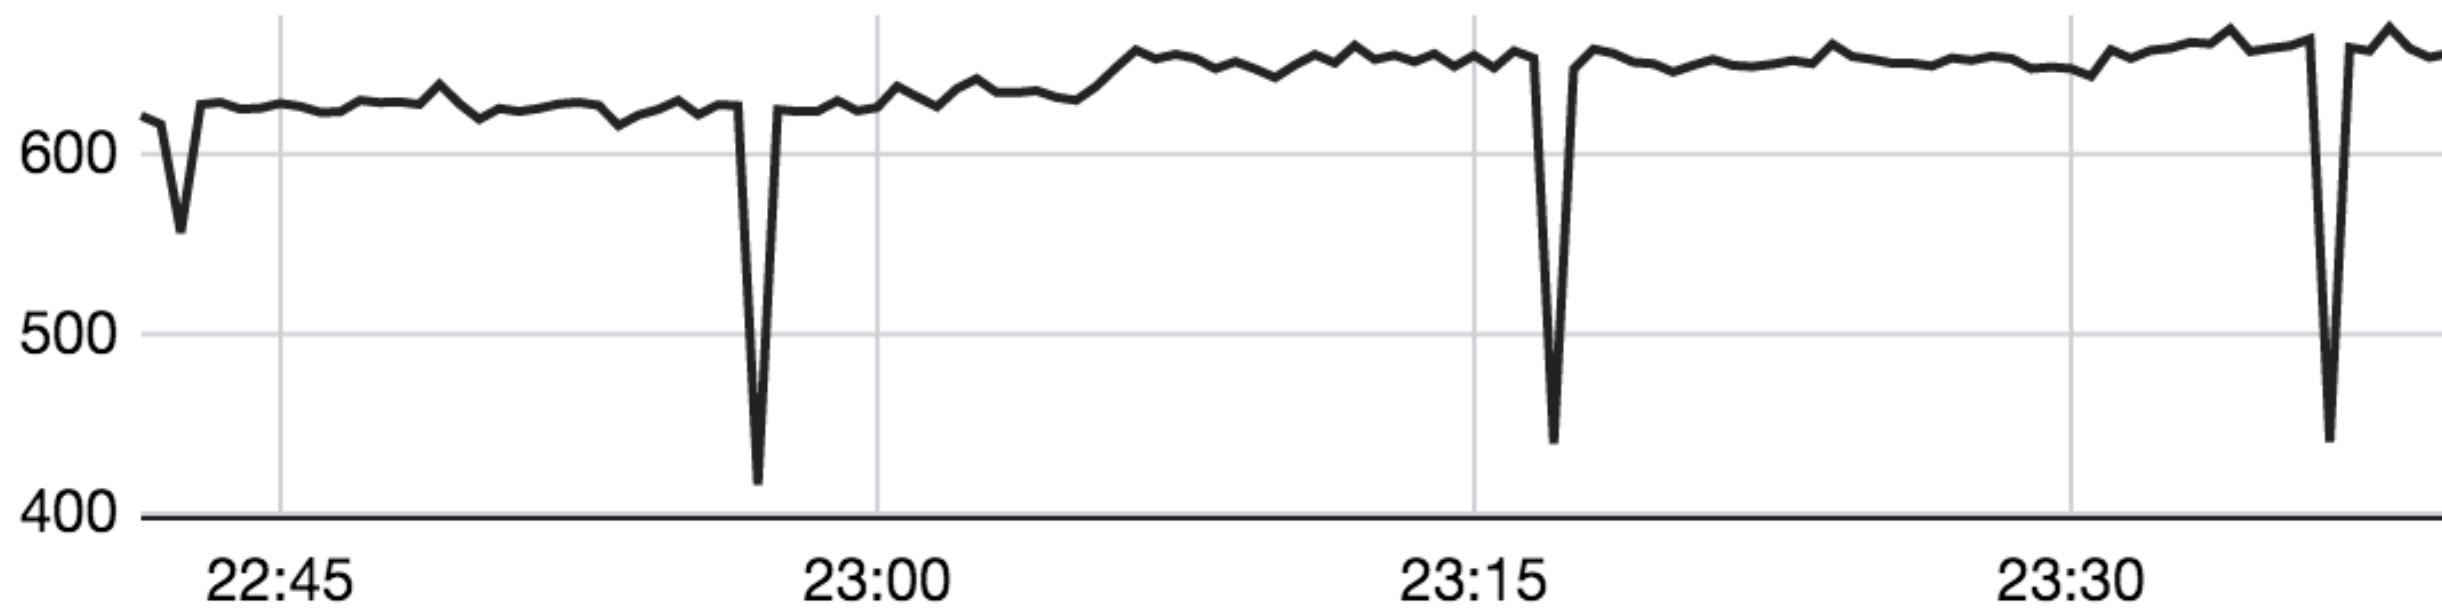
Buffering to protect Neo4j when indexing is too slow

Timeout due to the bug in the underlying implementation

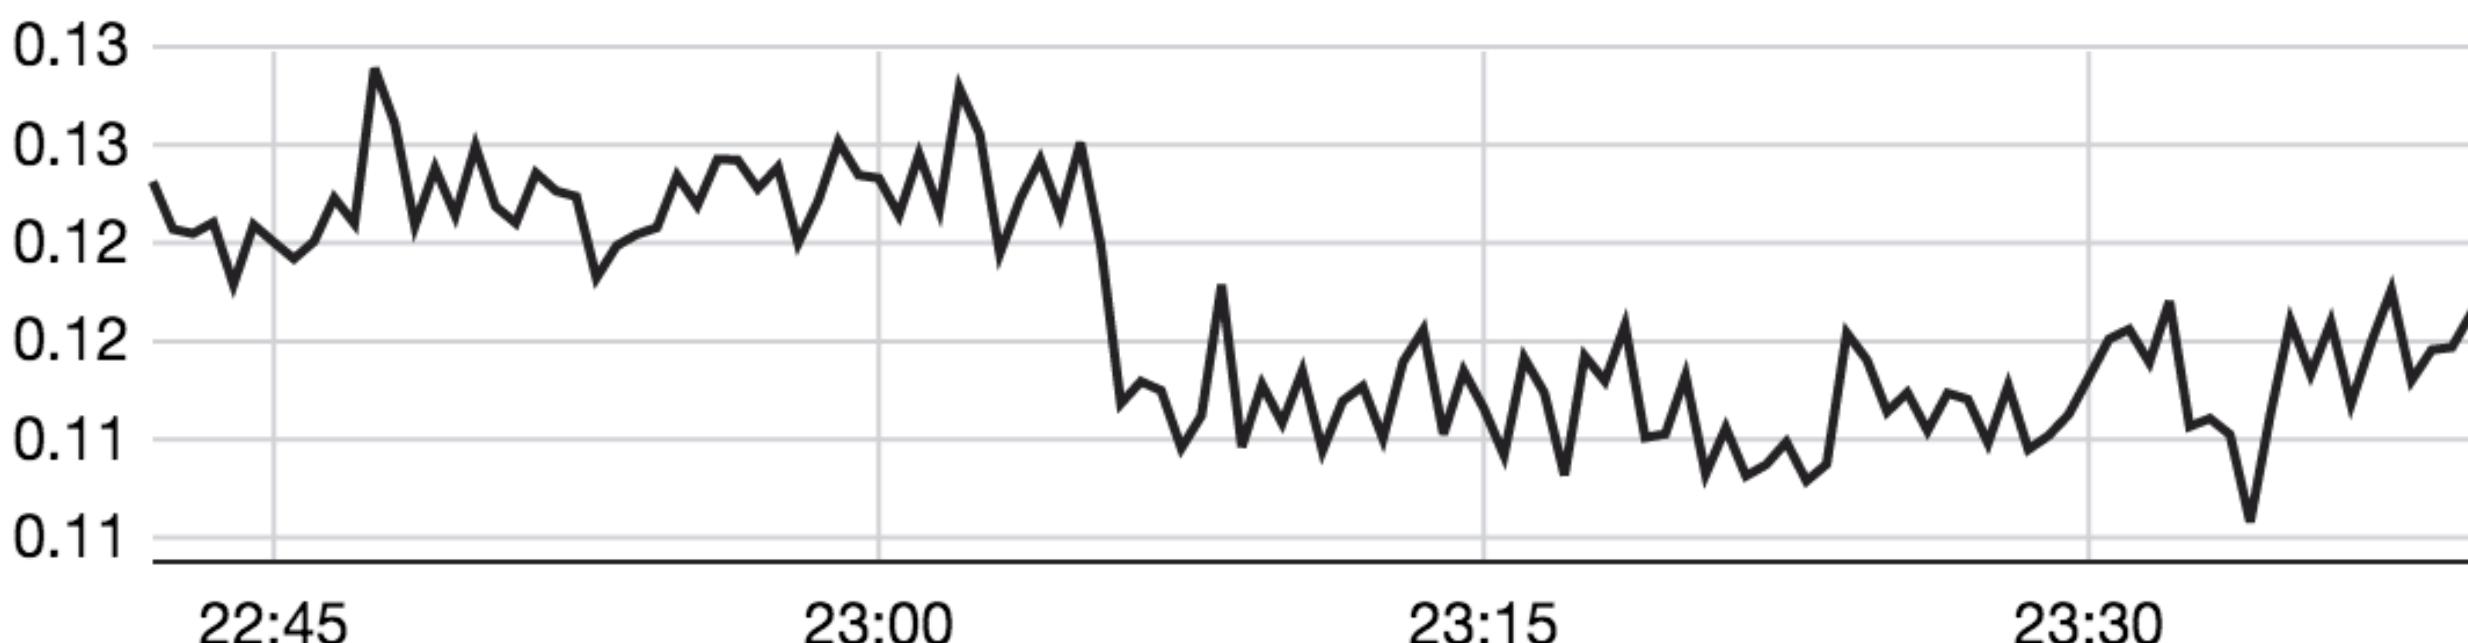


Bottleneck

Indexing Rate: 670.63 /s

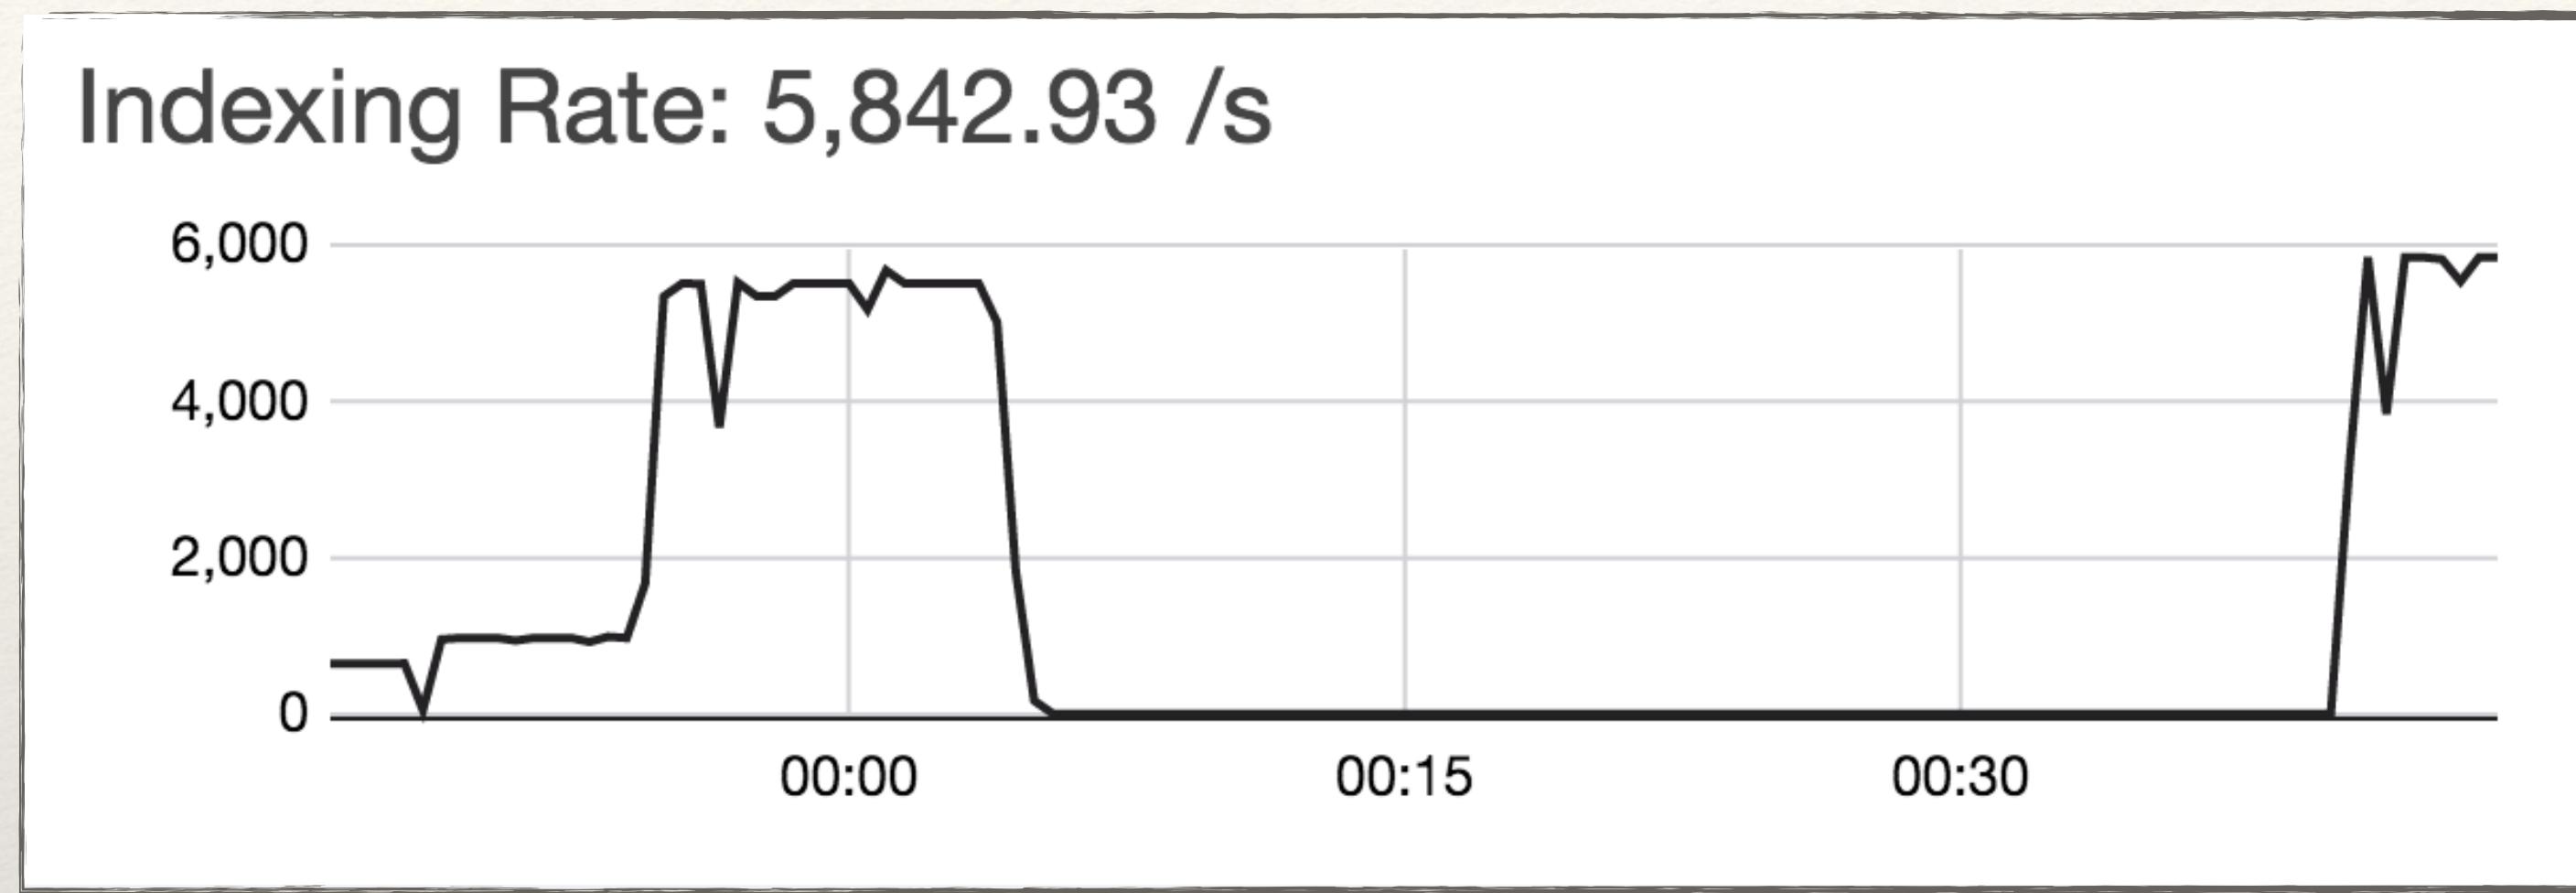


Indexing Latency: 0.11 ms



1.5 hour for 3 million subset, that's too long!

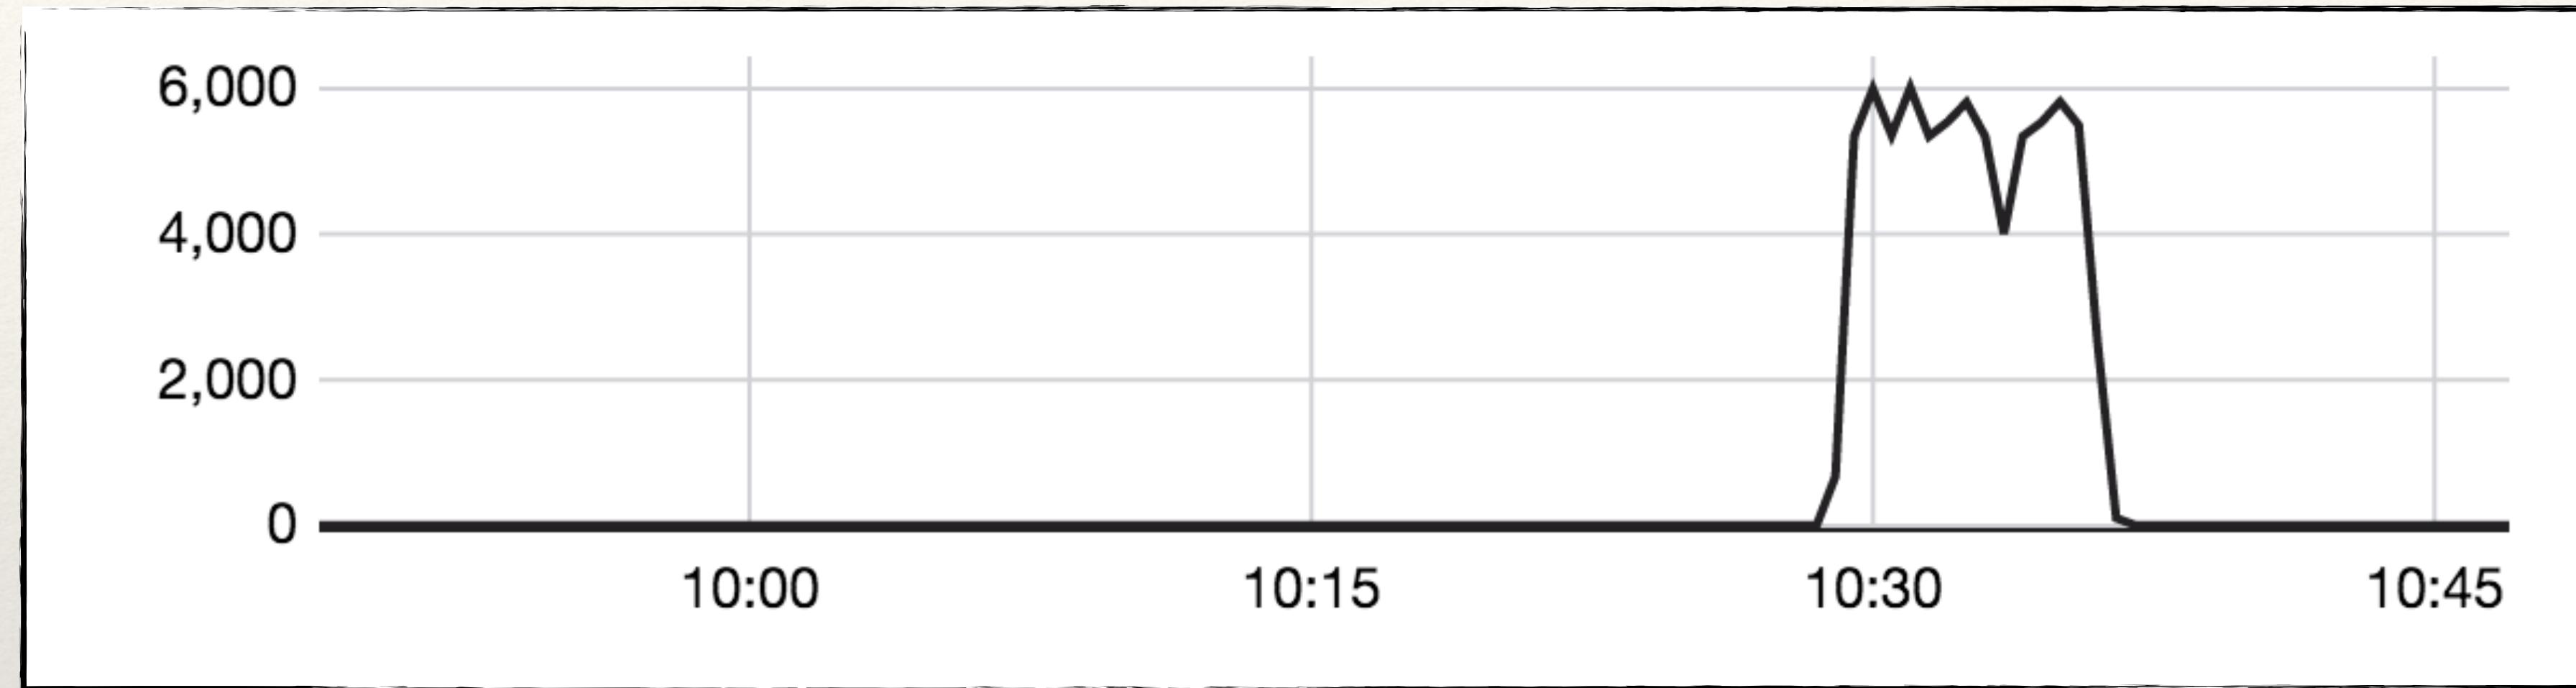
Bulk update with AkkaStream tuning



```
src
  .grouped(2000)
  .throttle(
    elements = 2,
    per = 6 second,
    maximumBurst = 2,
    mode = ThrottleMode.Shaping)
  .mapAsyncUnordered(parallelism = 1)(ids =>
    dao.bulkTag(ids, ...))
```

Bulk tagging,
few lines

Tagging Foo-Company



2,222,840 profiles matching criteria

~14 seconds until first batch is tagged

~7 minutes 11 seconds until all tagged

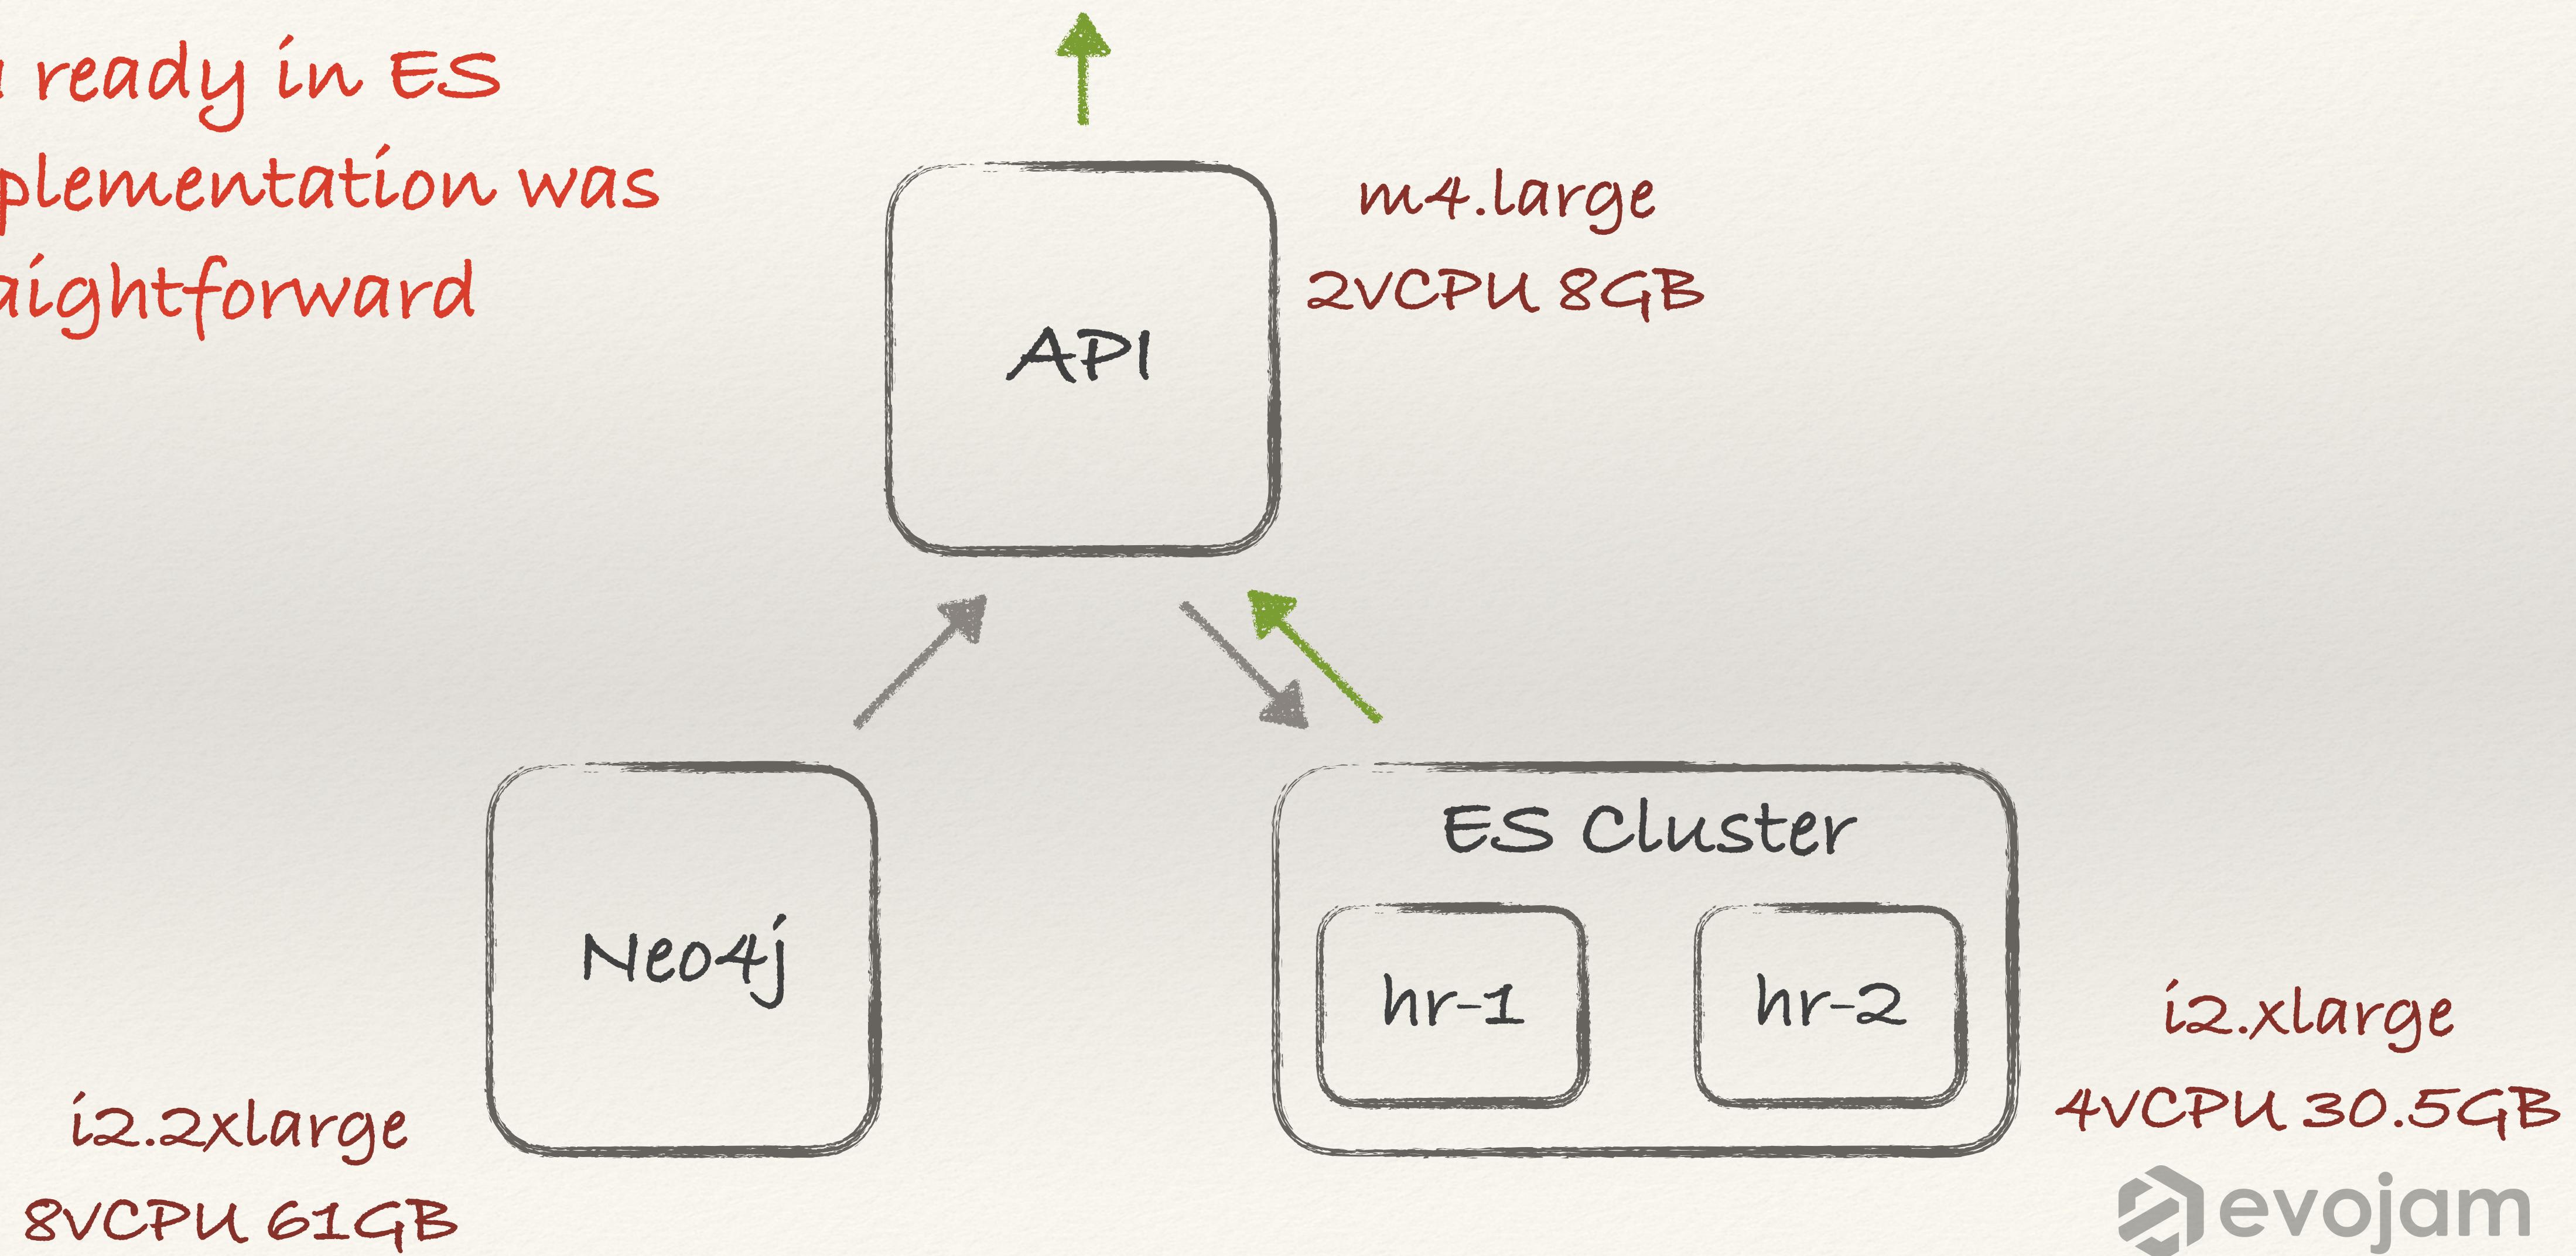
Sample subset :)

reference implementation - few hours

 evojam

Step #5 - Search

With data ready in ES
search implementation was
pretty straightforward



Step #5 - Search benchmark

GET /users?company=foo-company&phrase=John

Phrases based on
real names and
surnames, used
during profiles
enrichment

Fulltext search on 2 millions subset

2000 phrases for the benchmarking

Response JSON with 50 profiles

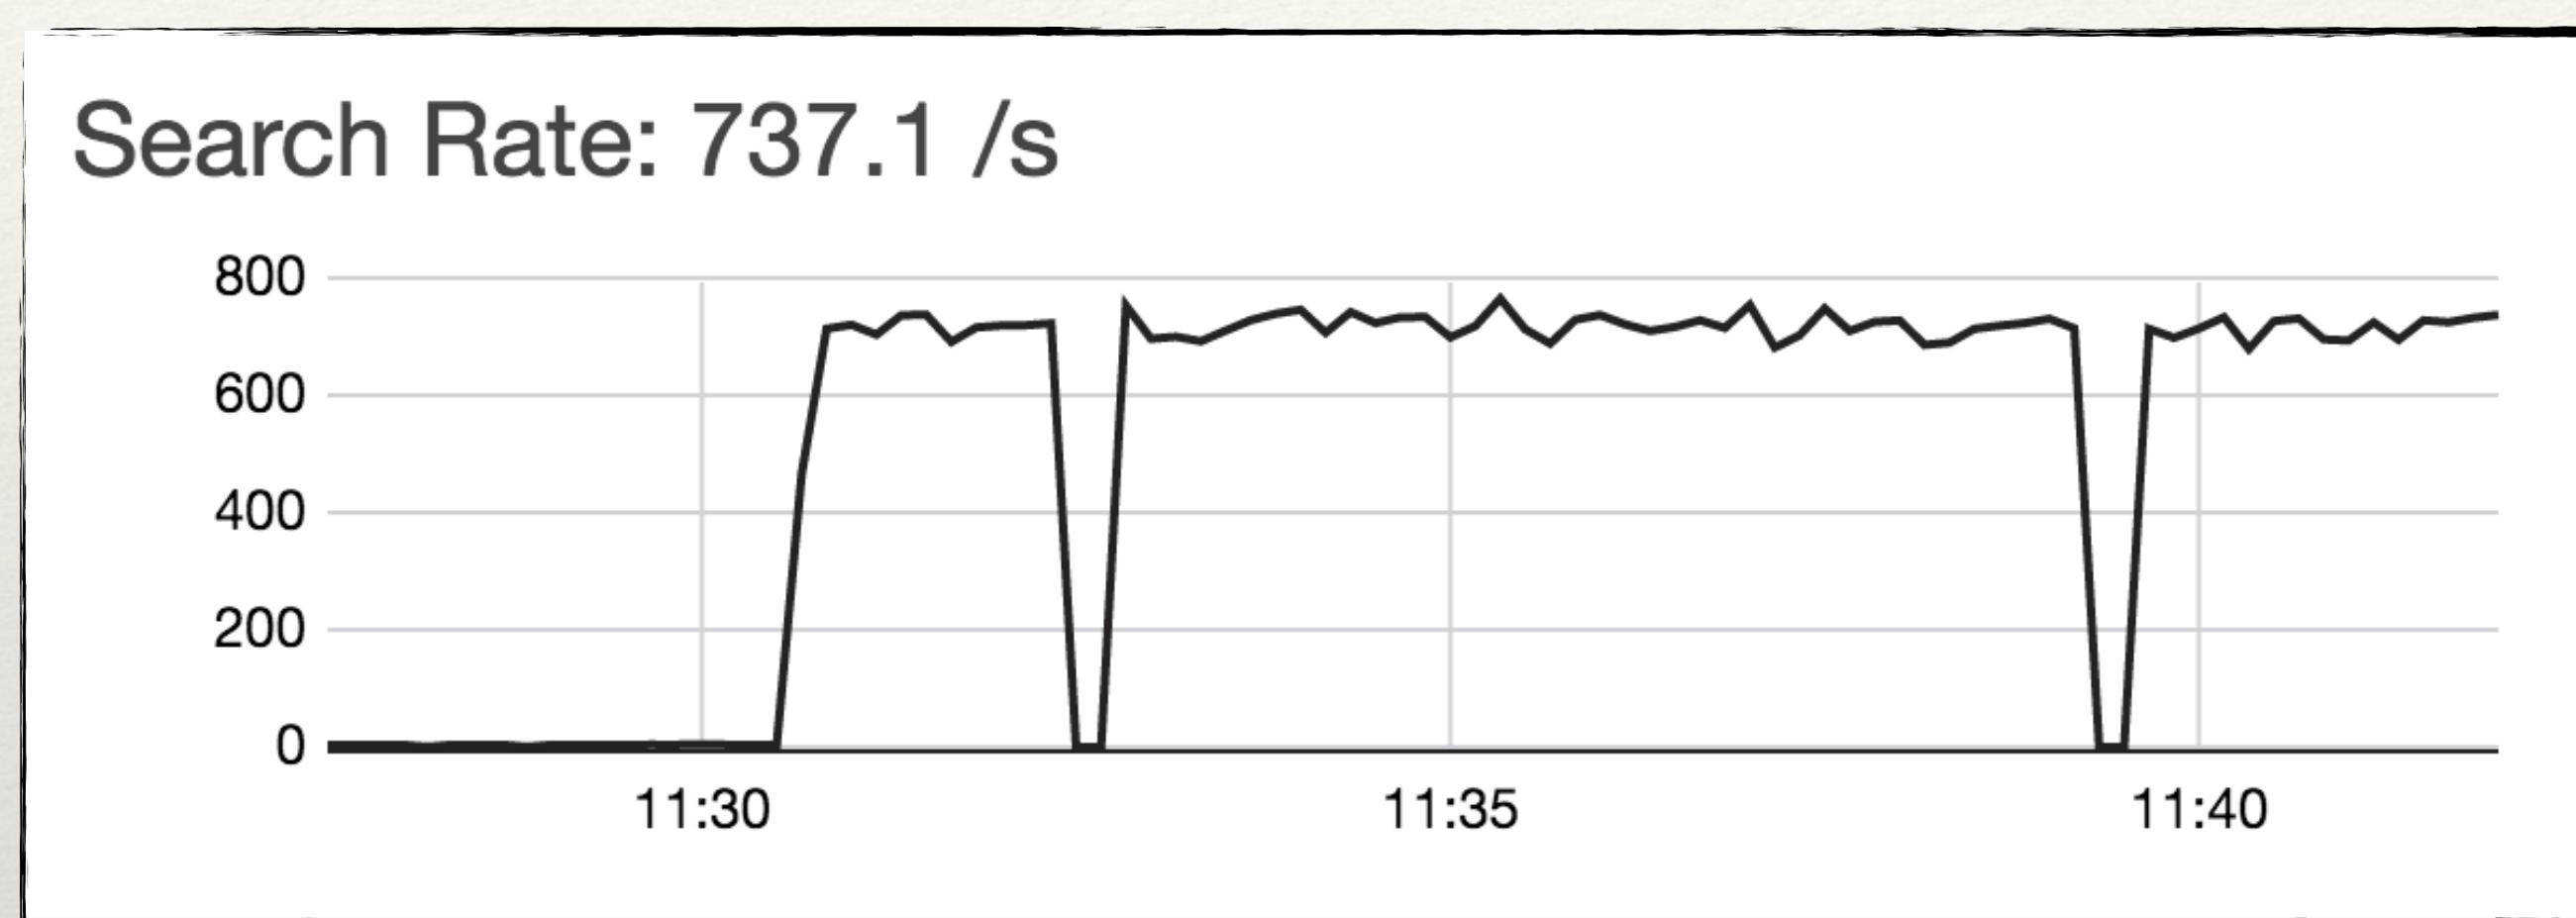
Searching in 750 millions database

Random
requests ordering

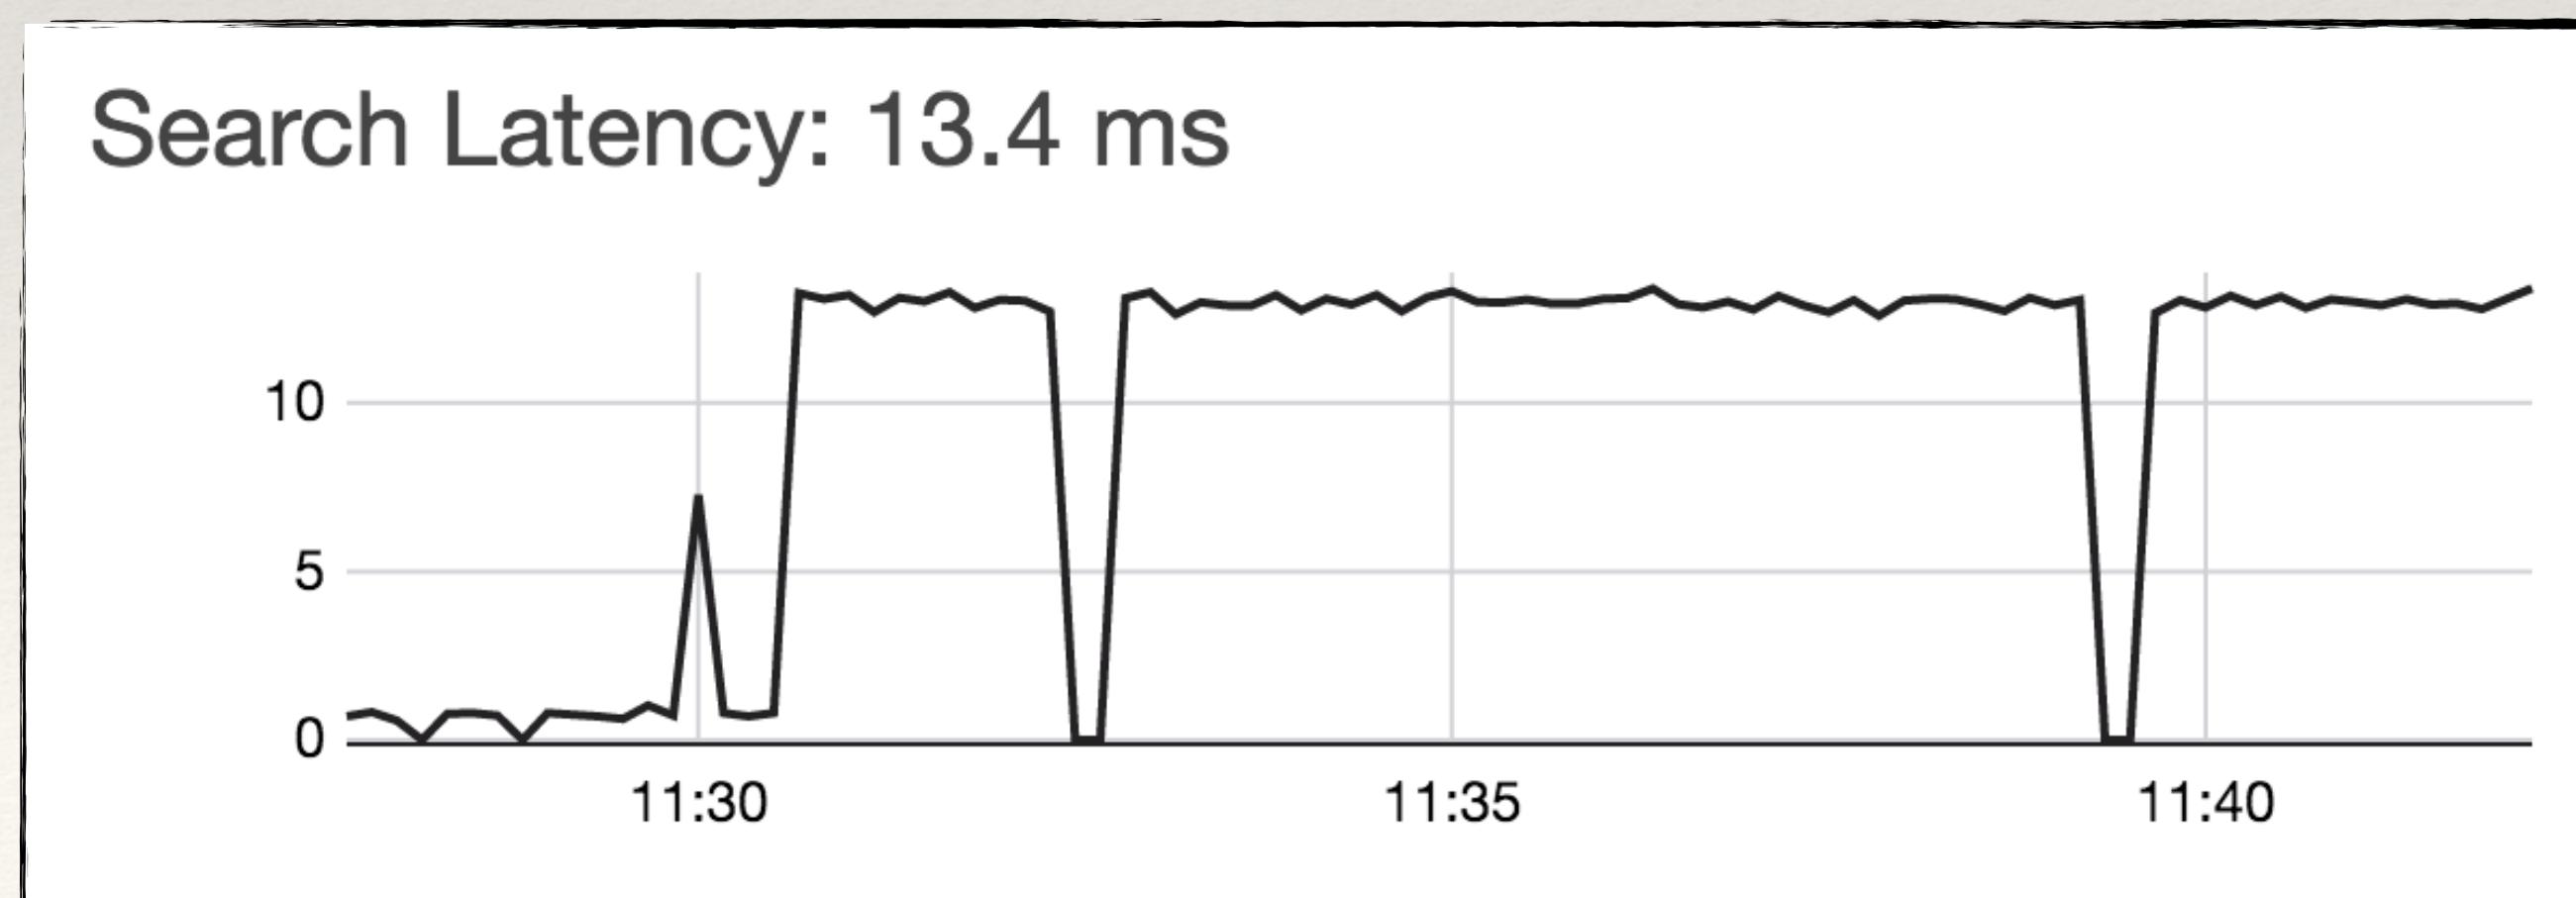
Step #5 - Search under siege

50 concurrent users

Response time ~ 0.14s



constant
search rate



constant
latency

Objective achieved

14 seconds until first batch is tagged

7 minutes until 2 millions ready

search response time from API ~0.14s

Summary

reference implementation

manual

few days

few days

~40 millions

no analytics

Tool ready for data scientists

We can implement core traversal modifications almost instantly

PoC

automatic

14 seconds

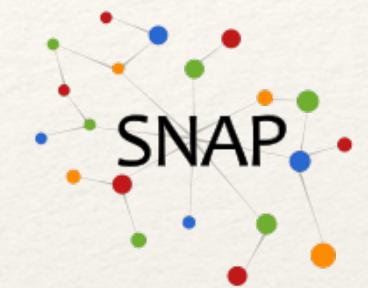
7 minutes

~750 millions

GraphX ready

Artur Bańkowski
artur@evojam.com
@abankowski

Summary



<http://neo4j.com/>

<https://www.elastic.co/>

<https://playframework.com/>

<http://doc.akka.io/docs/akka-stream-and-http-experimental/2.0/scala/stream-index.html>

Try at home!

<https://snap.stanford.edu/data/>

Sample graphs
ready to use

